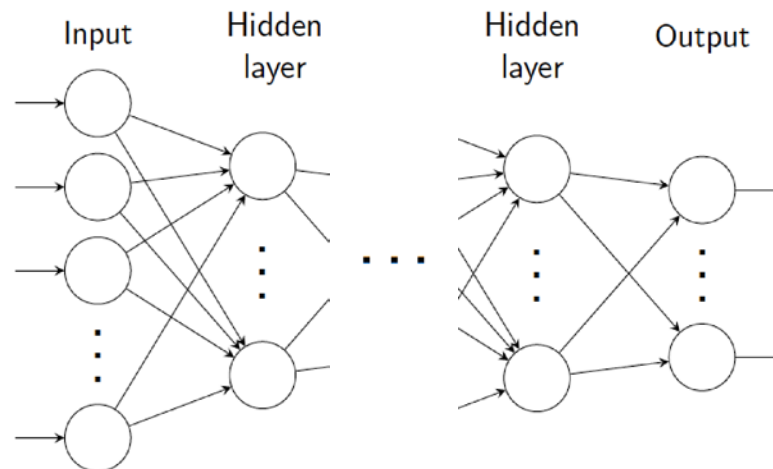


Training Neural Networks

6.036 Introduction to Machine Learning

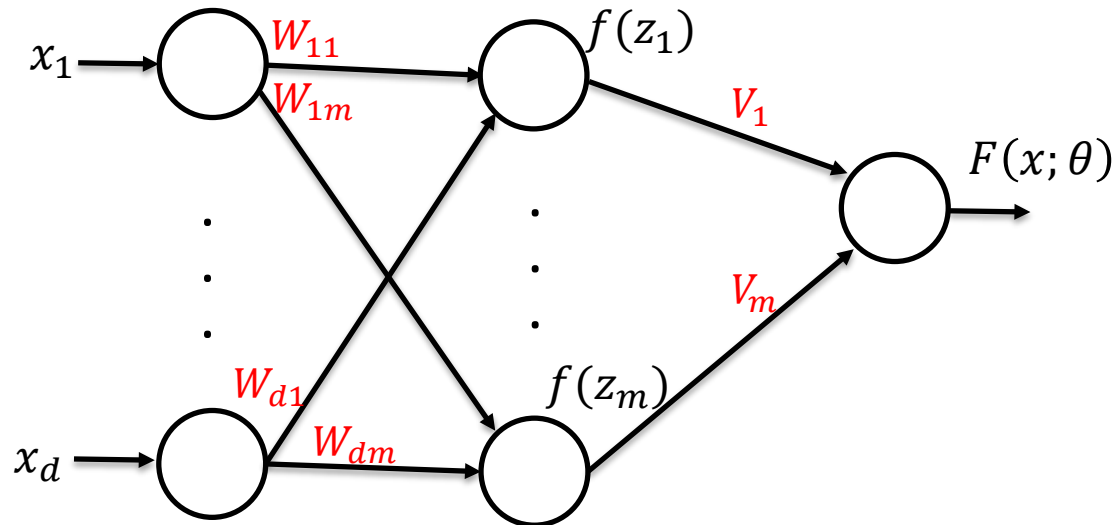
Feedforward Neural Networks

- Representation
 - Input, hidden layers, output
 - Parameters/weights
 - Activation functions
- Each layer computes some function of the previous layer
- Inputs mapped in a feed-forward fashion to output



Training Neural Networks

- Given a training dataset $S_n = \{(x^{(i)}, y^{(i)}), i = 1 \dots n\}$, estimate weights θ to minimize the average loss over the training examples
- $\theta = \{W_{ij}, W_{0j}, V_j, V_0\}$



Training using SGD

initialize network parameters θ

repeat (until some stopping criteria are met)

pick a random example \mathbf{t} from the training set

compute prediction: $F(\mathbf{x}^{(t)}; \theta)$

compute error/loss: $\text{Loss} \left(y^{(t)} F(\mathbf{x}^{(t)}; \theta) \right)$

compute gradient of the error: $\nabla_{\theta} \text{Loss} \left(y^{(t)} F(\mathbf{x}^{(t)}; \theta) \right)$

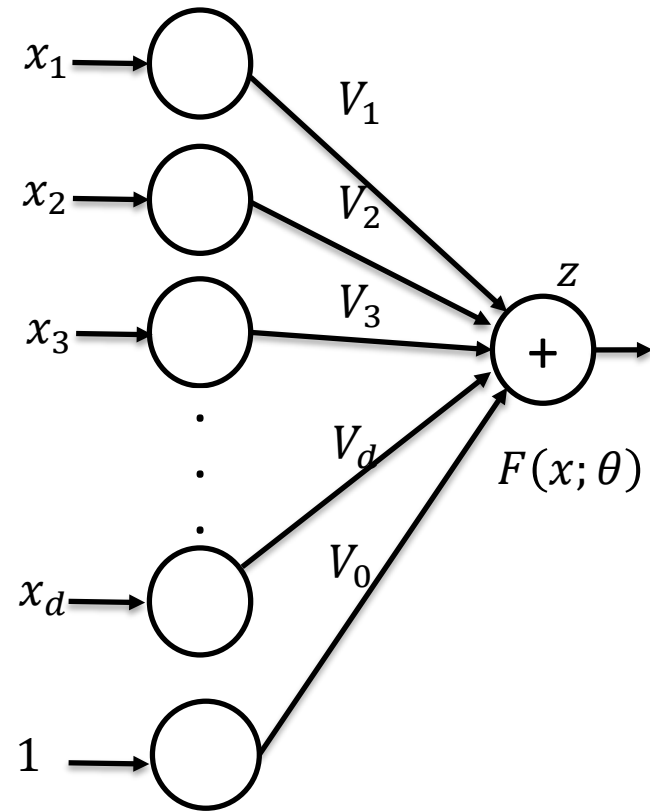
update weights: $\theta \leftarrow \theta - \eta \nabla_{\theta} \text{Loss} \left(y^{(t)} F(\mathbf{x}^{(t)}; \theta) \right)$

Gradients for Single-Layer Networks

- Computing gradient analytically

$$\begin{aligned}
 \frac{\partial}{\partial V_i} \text{Loss} \left(y^{(t)} F(x^{(t)}; \theta) \right) &= \\
 &= \frac{\partial}{\partial V_i} \text{Loss} \left(y^{(t)} z^{(t)} \right) \\
 &= \left[\frac{\partial}{\partial z^{(t)}} \text{Loss} \left(y^{(t)} z^{(t)} \right) \right] \left[\frac{\partial z^{(t)}}{\partial V_i} \right] = \\
 &= \left[\frac{\partial}{\partial z^{(t)}} \text{Loss} \left(y^{(t)} z^{(t)} \right) \right] \left[\frac{\partial \left(\sum_{j=1}^d x_j^{(t)} V_j + V_0 \right)}{\partial V_i} \right] = \\
 &= \begin{cases} -y^{(t)} & \text{if } \text{Loss}(y^{(t)} z^{(t)}) > 0 \\ 0 & \text{otherwise} \end{cases} \begin{bmatrix} x_i^{(t)} \end{bmatrix}
 \end{aligned}$$

Chain Rule: $\frac{dy}{dx} = \frac{dy}{du} \frac{du}{dx}$



$$z^{(t)} = \sum_{j=1}^d x_j^{(t)} V_j + V_0$$

$$F(x^{(t)}; \theta) = f(z^{(t)}) = z^{(t)}$$

Training using SGD

initialize network parameters $\theta = \{V_0, \dots, V_d\}$

repeat (until some stopping criteria are met)

pick a random example \mathbf{t} from the training set

compute prediction: $F(\mathbf{x}^{(t)}; \theta)$

compute error/loss: $\text{Loss} \left(y^{(t)} F(\mathbf{x}^{(t)}; \theta) \right)$

compute gradient of the error: $\nabla_{\theta} \text{Loss} \left(y^{(t)} F(\mathbf{x}^{(t)}; \theta) \right)$

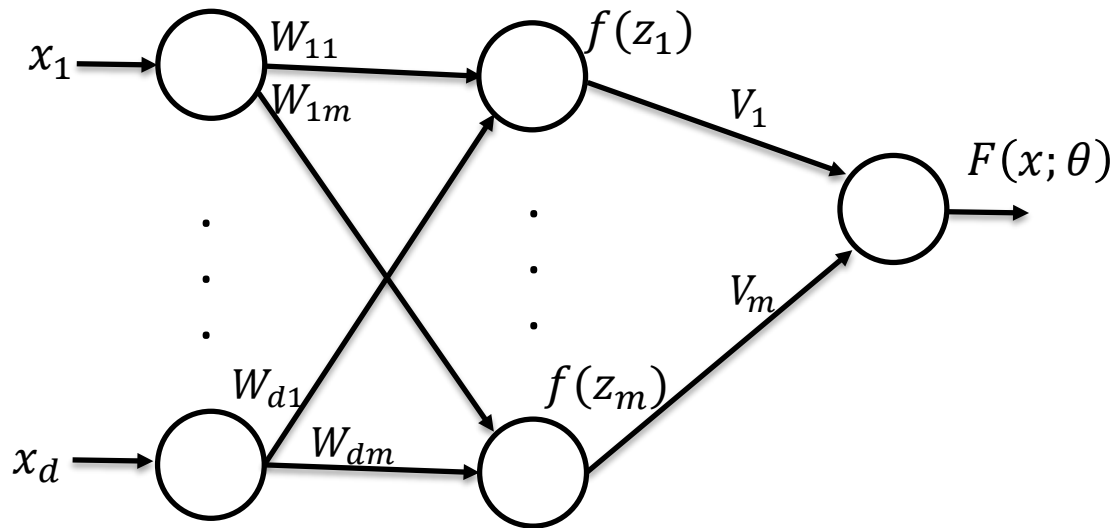
update weights: $\theta \leftarrow \theta - \eta \nabla_{\theta} \text{Loss} \left(y^{(t)} F(\mathbf{x}^{(t)}; \theta) \right)$

$$V_i \leftarrow V_i + \eta x_i^{(t)} y^{(t)}, i = 1 \dots d$$

$$V_0 \leftarrow V_0 + \eta y^{(t)}$$

SGD for 2-Layer Neural Network

Parameters $\theta = \{W_{ij}, W_{0j}\} \& \{V_j, V_0\}$



$$z_j = \sum_{i=1}^d x_i W_{ij} + W_{0j}$$

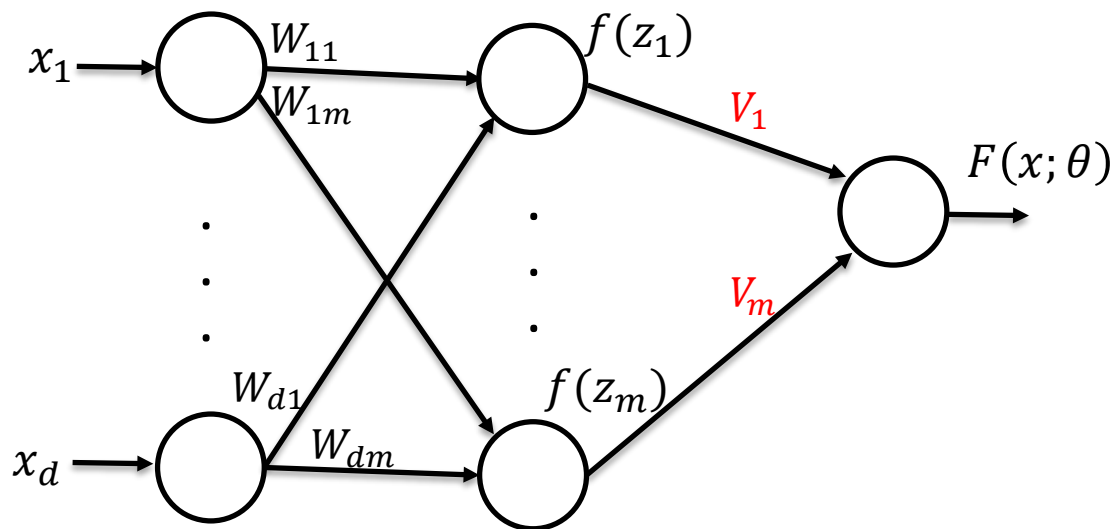
$$f(z_j) = \max\{0, z_j\}$$

$$z = \sum_{j=1}^m f(z_j) V_j + V_0$$

$$F(x; \theta) = z$$

Updates for V_j

Parameters $\theta = \{W_{ij}, W_{0j}\} \& \{V_j, V_0\}$



$$z_j = \sum_{i=1}^d x_i W_{ij} + W_{0j}$$

$$f(z_j) = \max\{0, z_j\}$$

$$z = \sum_{j=1}^m f(z_j) V_j + V_0$$

$$F(x; \theta) = z$$

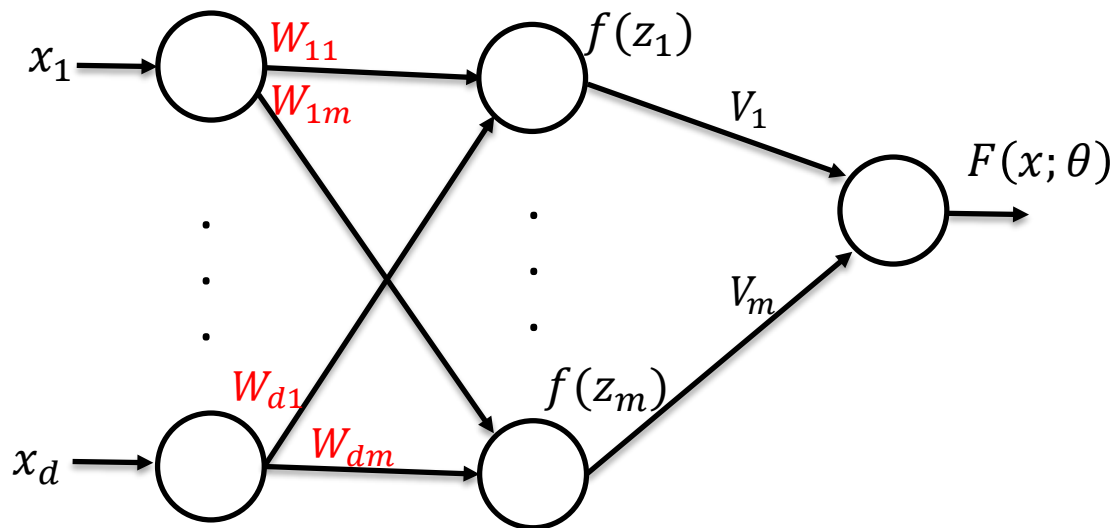
Updates for V_j are the same as for the single-layer network

Replace $x_i^{(t)}$ with $f(z_j^{(t)})$

$$V_j \leftarrow V_j + \eta_k y^{(t)} f(z_j^{(t)}), \quad j = 1, \dots, m$$

Updates for W_{ij}

Parameters $\theta = \{W_{ij}, W_{0j}\} \& \{V_j, V_0\}$



$$z_j = \sum_{i=1}^d x_i W_{ij} + W_{0j}$$

$$f(z_j) = \max\{0, z_j\}$$

$$z = \sum_{j=1}^m f(z_j) V_j + V_0$$

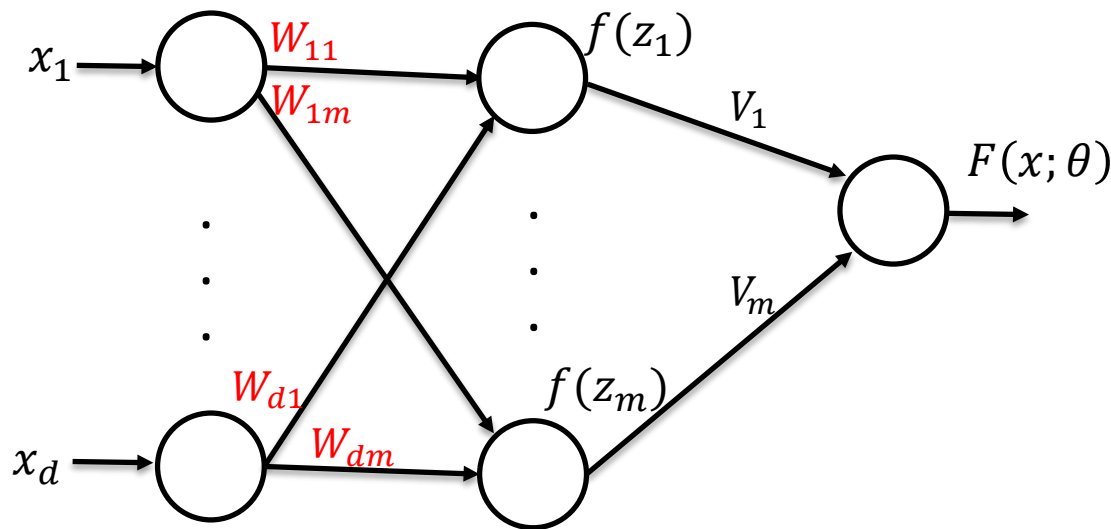
$$F(x; \theta) = z$$

To compute $\frac{\partial}{\partial V_i} \text{Loss} \left(y^{(t)} F(x^{(t)}; \theta) \right)$, we could start by writing

$\text{Loss} \left(y^{(t)} F(x^{(t)}; \theta) \right)$ as a function of the network parameters θ . And then compute the partial derivatives ... Instead, we can use the chain rule to derive a compact algorithm: **back-propagation**

Updates for W_{ij}

Parameters $\theta = \{W_{ij}, W_{0j}\} \& \{V_j, V_0\}$



$$z_j = \sum_{i=1}^d x_i W_{ij} + W_{0j}$$

$$f(z_j) = \max\{0, z_j\}$$

$$z = \sum_{j=1}^m f(z_j) V_j + V_0$$

$$F(x; \theta) = z$$

$$\begin{aligned} \frac{\partial}{\partial W_{ij}} \text{Loss}(y^{(t)} z^{(t)}) &= \left[\frac{\partial}{\partial z^{(t)}} \text{Loss}(y^{(t)} z^{(t)}) \right] \left[\frac{\partial z^{(t)}}{\partial f(z_j^{(t)})} \right] \left[\frac{\partial f(z_j^{(t)})}{\partial z_j^{(t)}} \right] \left[\frac{\partial z_j^{(t)}}{\partial W_{ij}} \right] \\ &= \begin{bmatrix} -y^{(t)} & \text{if } \text{Loss}(y^{(t)} z^{(t)}) > 0 \\ 0 & \text{otherwise} \end{bmatrix} [V_j] \llbracket z_j^{(t)} > 0 \rrbracket [x_i] \end{aligned}$$

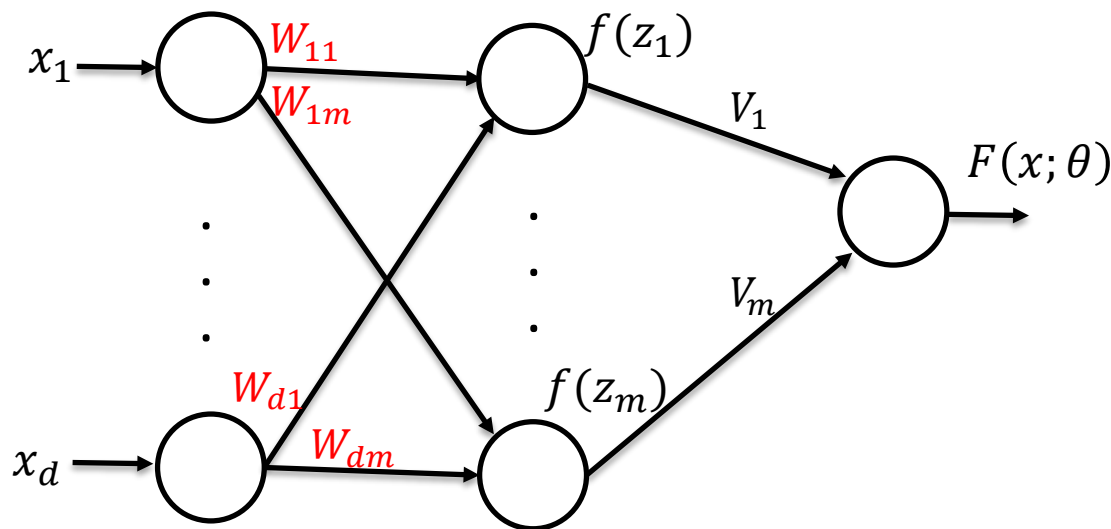
Back-propagation

- The process of propagating the gradients backwards towards the input layer is called **back-propagation**
- Back-propagation is based on applying the chain rule of derivatives back through the model
- Back-propagation can be applied the same way to compute gradients in networks with multiple hidden layers
- Computing derivatives can be reused between layers

$$\begin{aligned}\frac{\partial}{\partial W_{ij}} \text{Loss}(y^{(t)} z^{(t)}) &= \left[\frac{\partial}{\partial z^{(t)}} \text{Loss}(y^{(t)} z^{(t)}) \right] \left[\frac{\partial z^{(t)}}{\partial f(z_j^{(t)})} \right] \left[\frac{\partial f(z_j^{(t)})}{\partial z_j^{(t)}} \right] \left[\frac{\partial z_j^{(t)}}{\partial W_{ij}} \right] \\ &= \begin{bmatrix} -y^{(t)} & \text{if } \text{Loss}(y^{(t)} z^{(t)}) > 0 \\ 0 & \text{otherwise} \end{bmatrix} [V_j] \llbracket z_j^{(t)} > 0 \rrbracket [x_i]\end{aligned}$$

Updates for W_{ij}

Parameters $\theta = \{W_{ij}, W_{0j}\} \& \{V_j, V_0\}$



$$z_j = \sum_{i=1}^d x_i W_{ij} + W_{0j}$$

$$f(z_j) = \max\{0, z_j\}$$

$$z = \sum_{j=1}^m f(z_j) V_j + V_0$$

$$F(x; \theta) = z$$

$$W_{ij} \leftarrow W_{ij} + \eta_k x_i^{(t)} \llbracket z_j^{(t)} > 0 \rrbracket V_j y^{(t)}, \quad i = 1, \dots, d, \quad j = 1, \dots, m$$

Initialization

- What happens when all weights are initialized to 0?

$$V_j \leftarrow V_j + \eta_k y^{(t)} f(z_j^{(t)}), \quad j = 1, \dots, m$$

$$W_{ij} \leftarrow W_{ij} + \eta_k x_i^{(t)} \llbracket z_j^{(t)} > 0 \rrbracket V_j y^{(t)}, \quad i = 1, \dots, d, \quad j = 1, \dots, m$$

Initialization

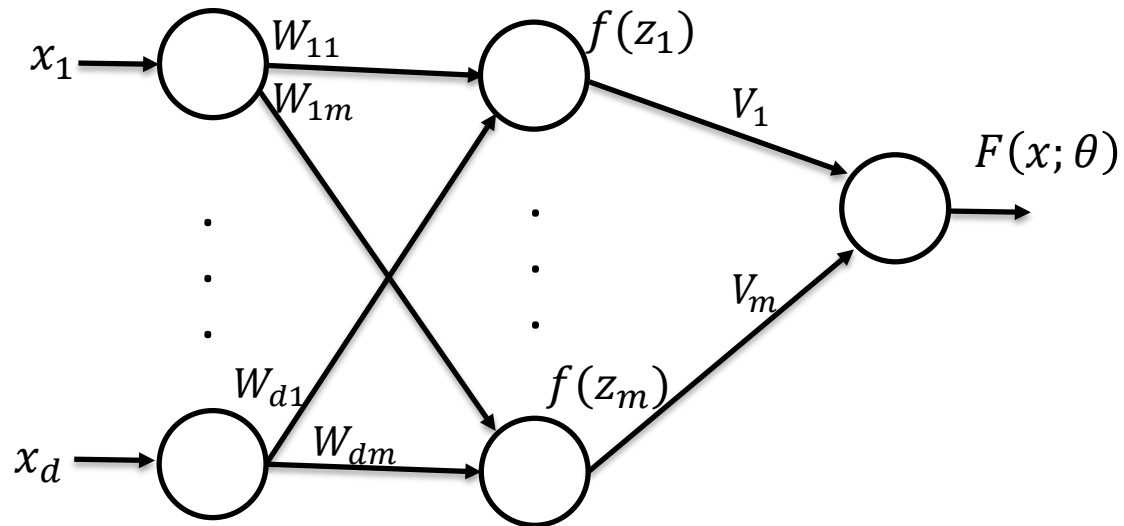
- What happens when all weights are initialized to 0?

$$V_j \leftarrow V_j$$

$$W_{ij} \leftarrow W_{ij}$$

Initialization

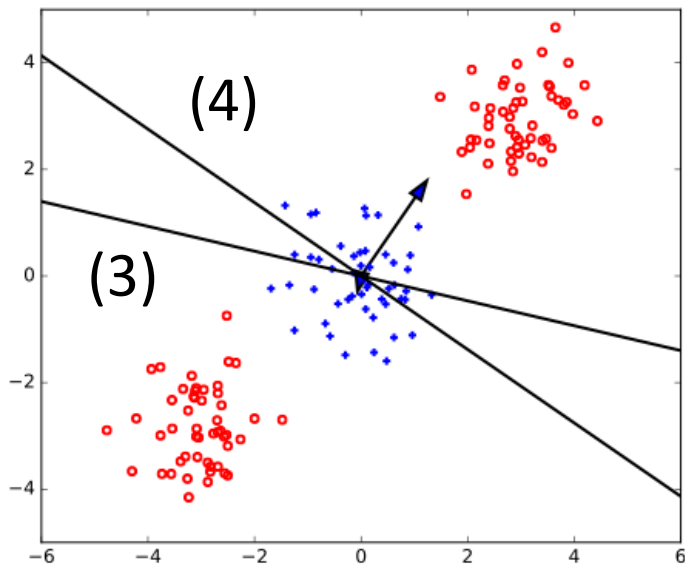
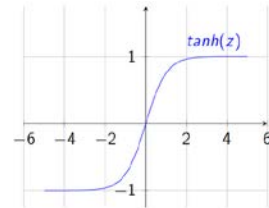
- Typically random
 - Zero mean and variance $1/d^2$



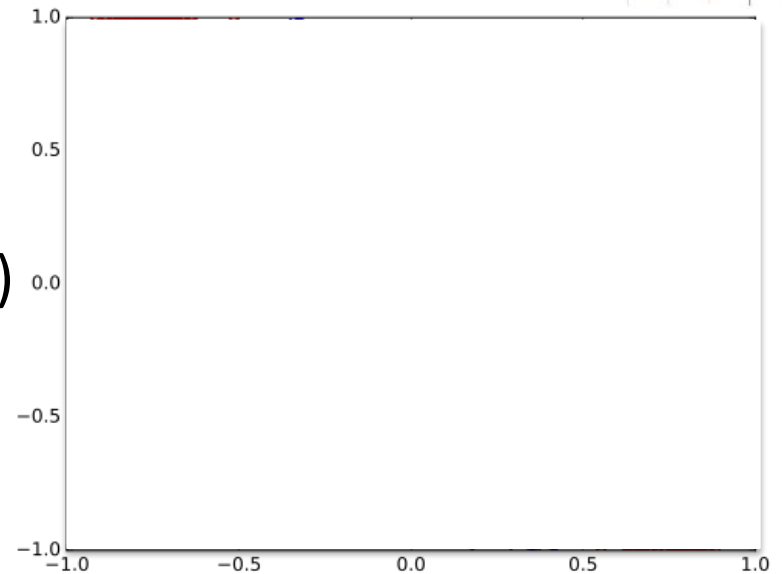
2 Hidden Units

- Randomly initialized weights (zero offset)

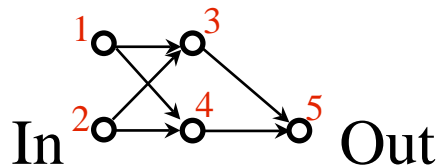
tanh activation



(4)



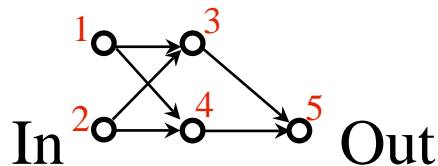
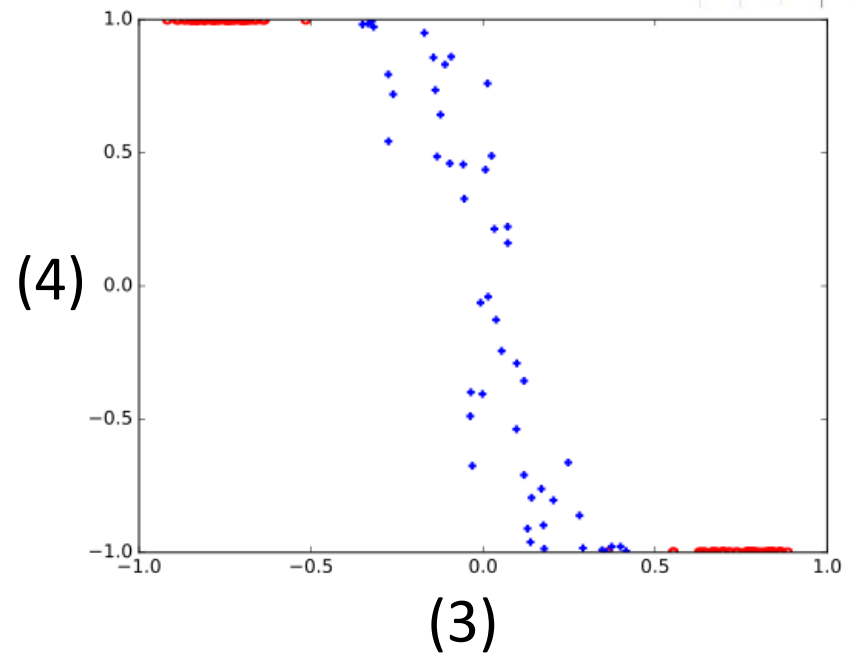
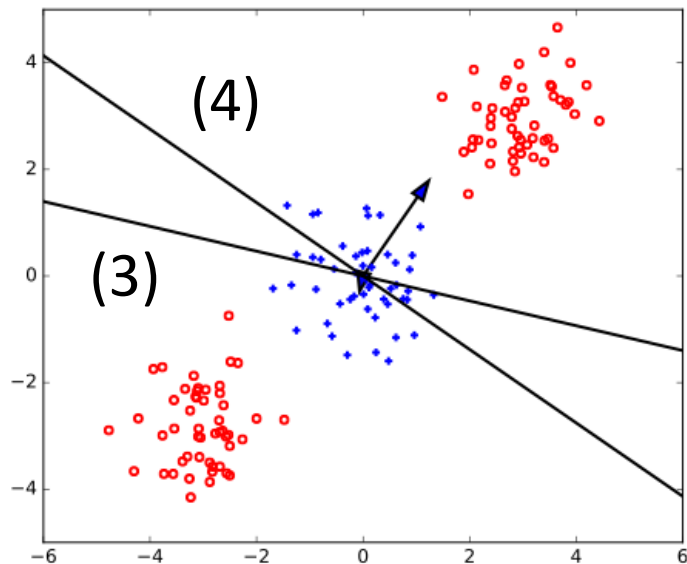
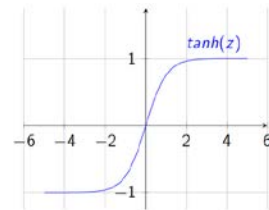
(3)



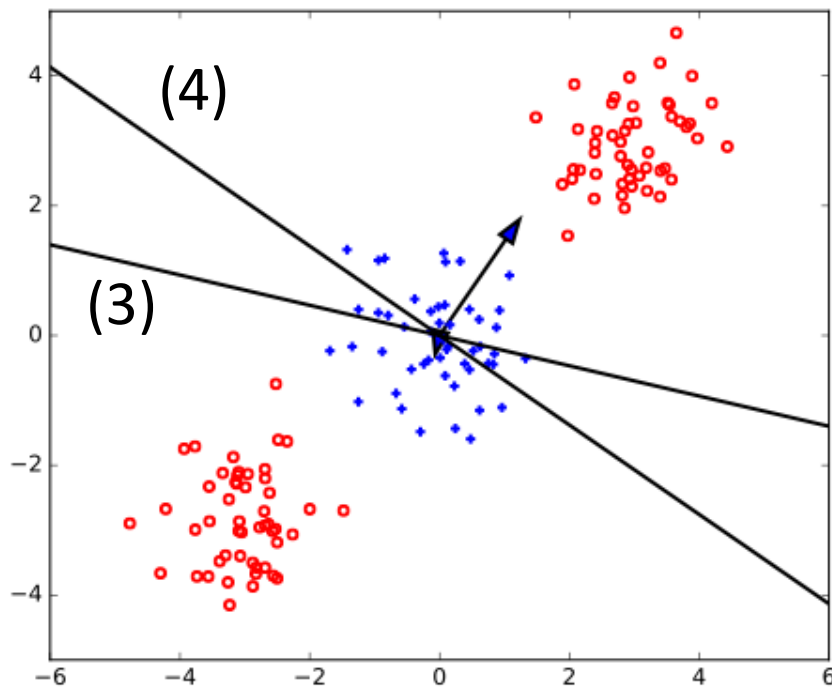
2 Hidden Units

- Randomly initialized weights (zero offset)

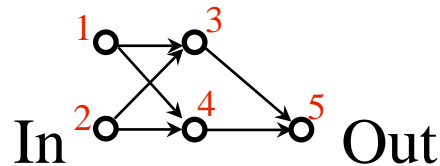
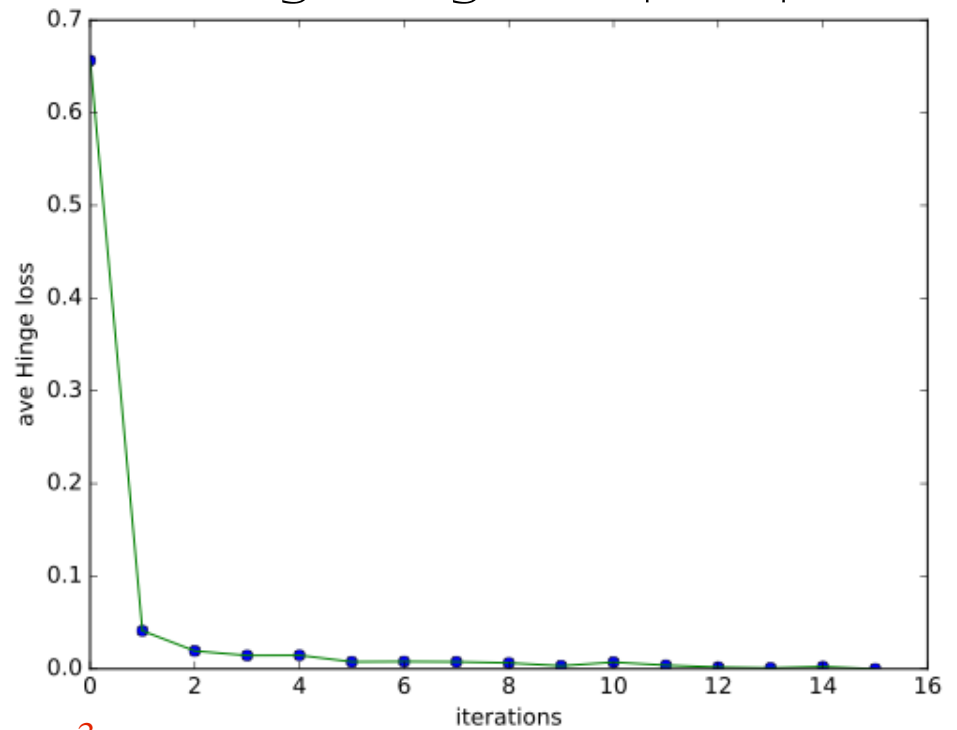
tanh activation



2 Hidden Units: Training

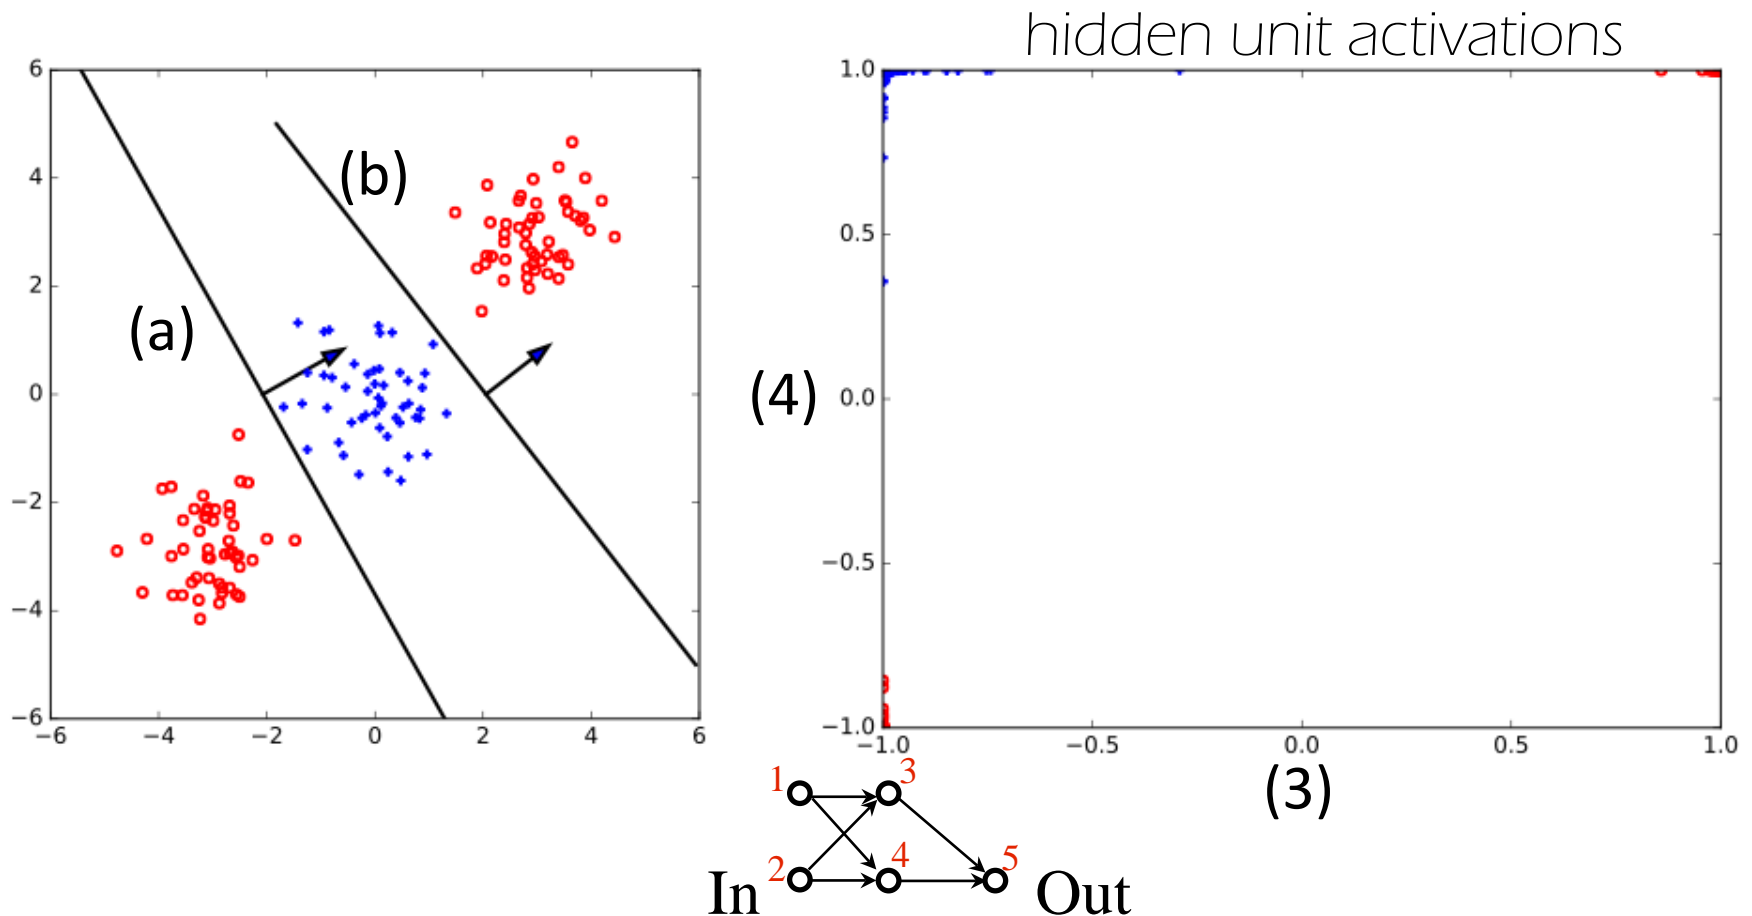


Average hinge loss per epoch



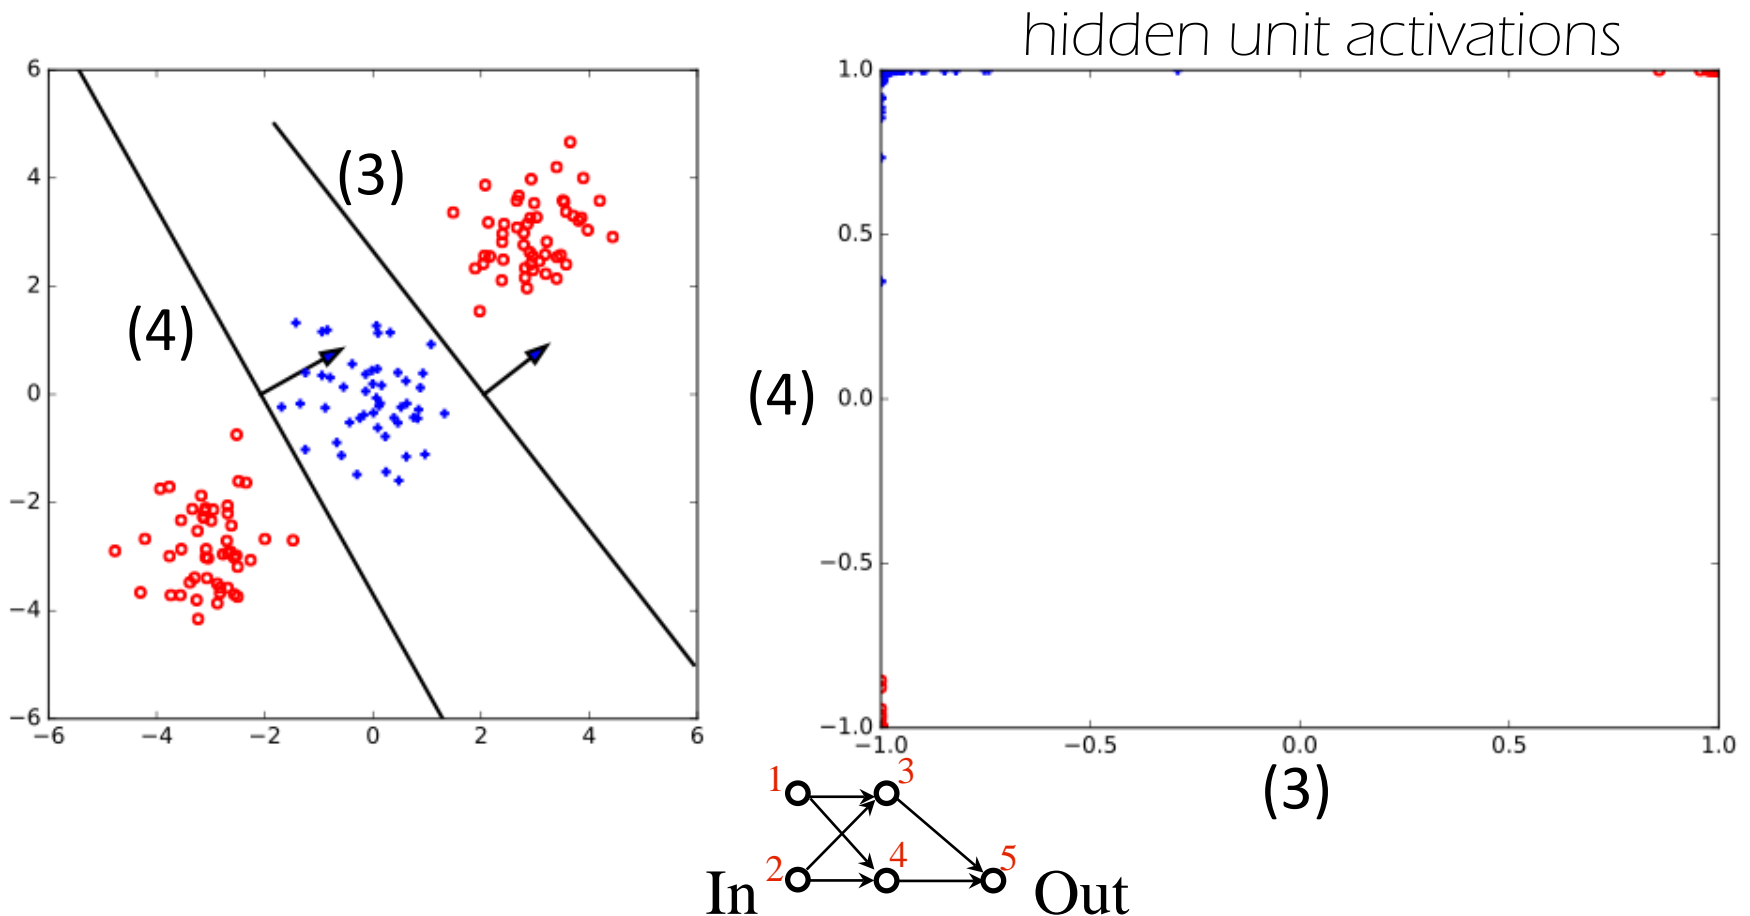
2 Hidden Units: Training

- After ~ 10 passes through the data



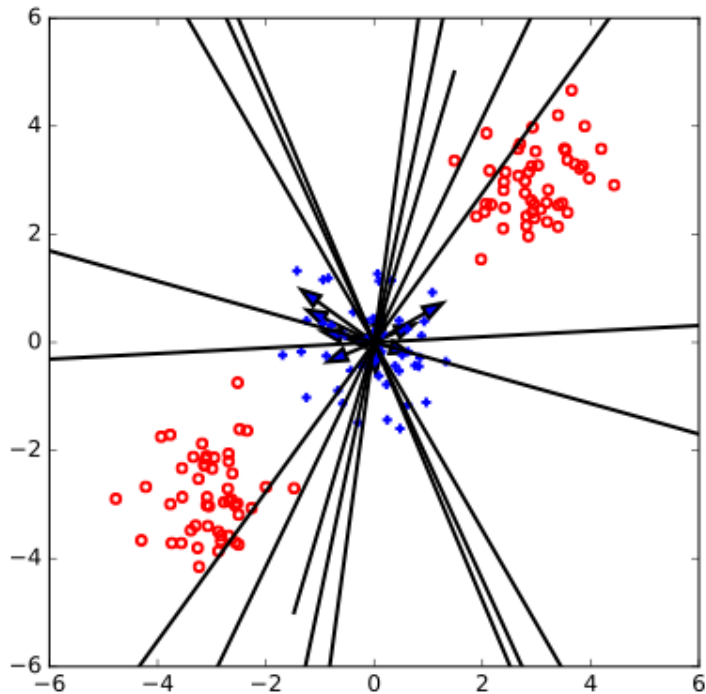
2 Hidden Units: Training

- After ~ 10 passes through the data



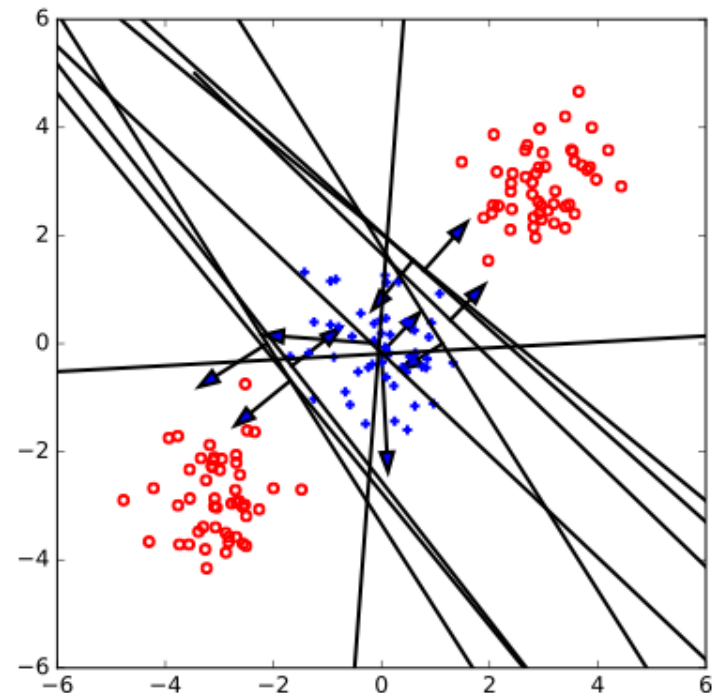
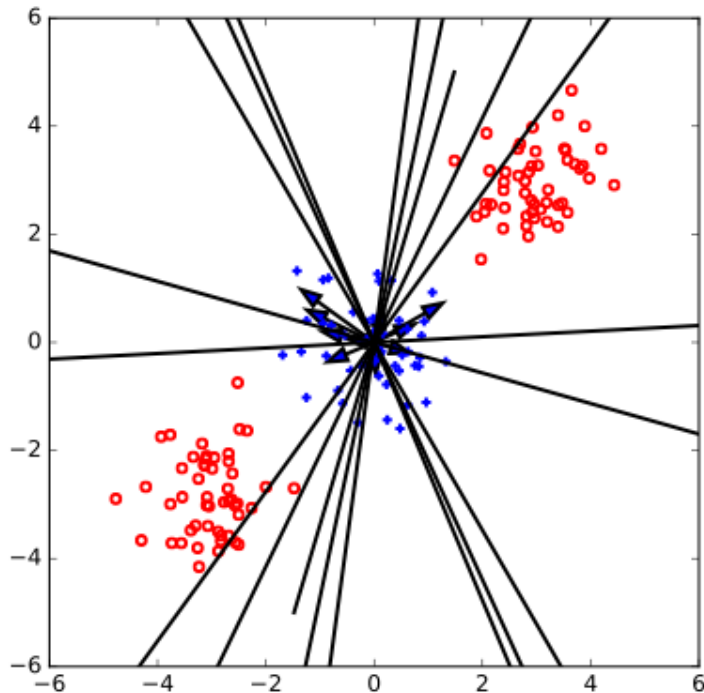
10 Hidden Units

- Randomly initialized weights (zero offset) for the hidden units



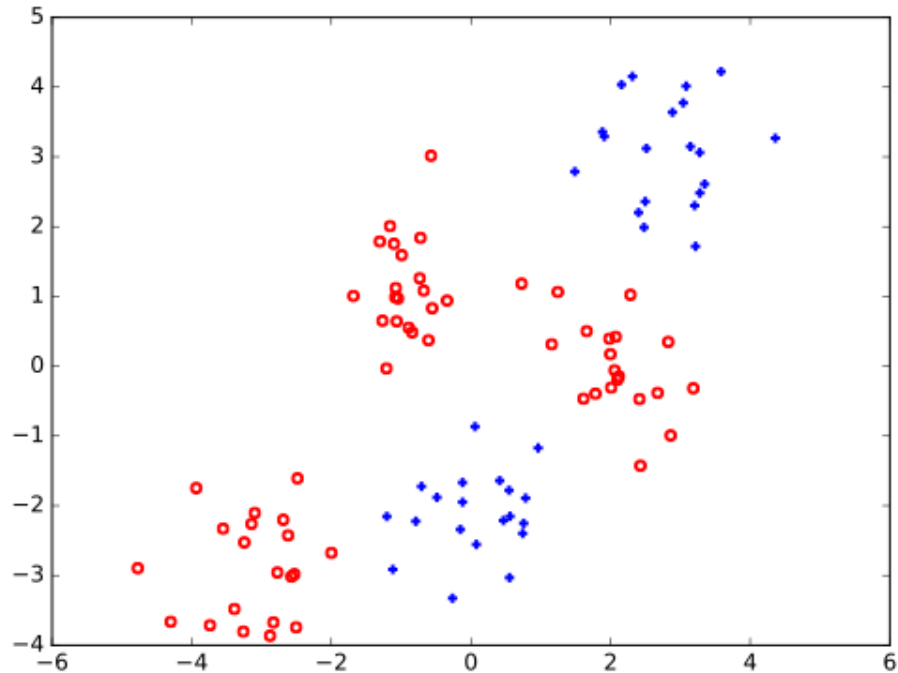
10 Hidden Units

- After ~ 10 epochs the hidden units are arranged in a manner sufficient for the task (but not otherwise perfect)



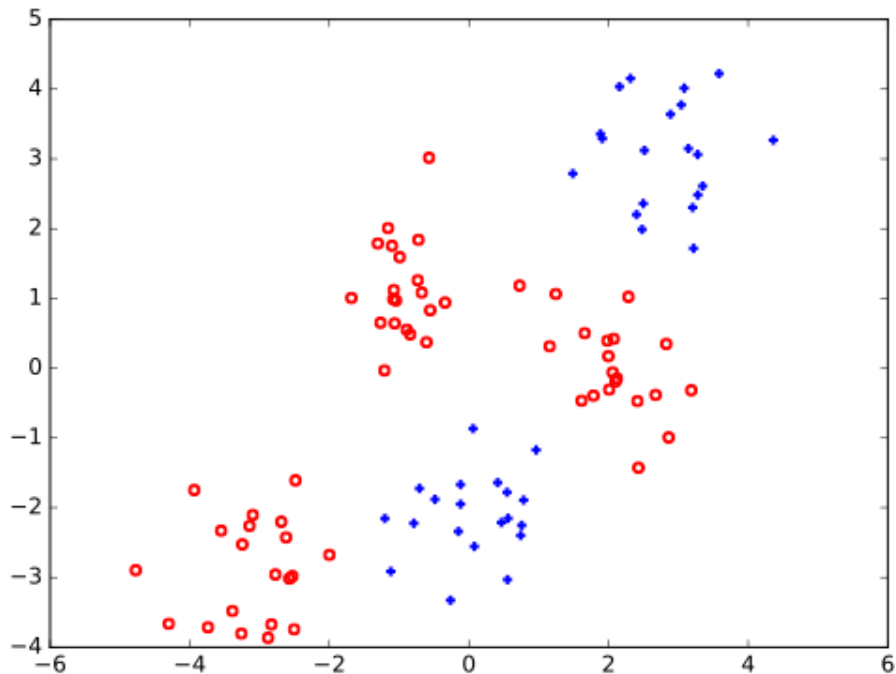
Decisions (and a harder task)

- 2 hidden units can no longer solve this task

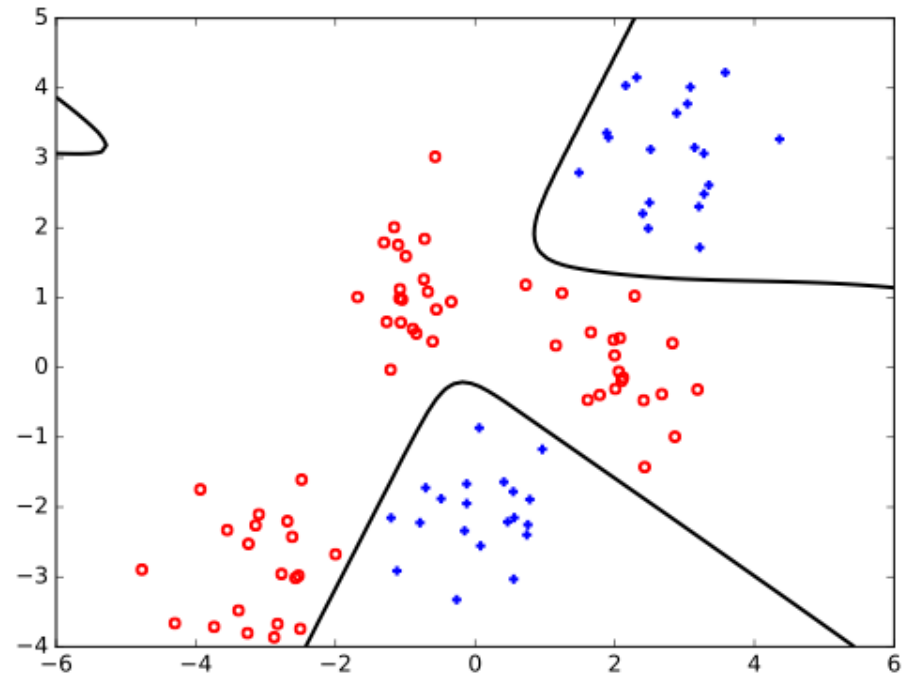


Decisions (and a harder task)

- 2 hidden units can no longer solve this task

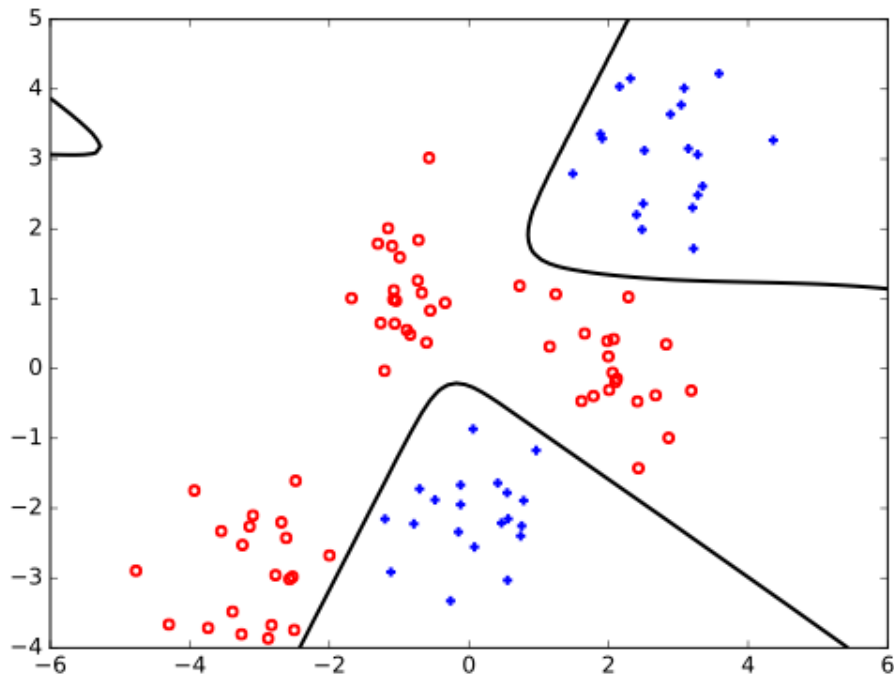


10 hidden units

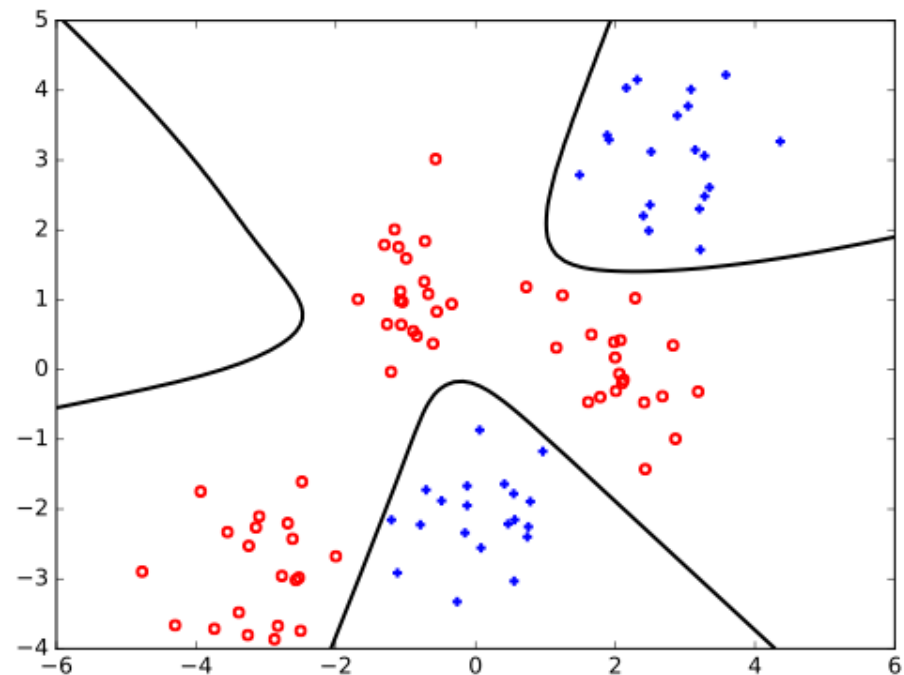


Decisions (and a harder task)

10 hidden units

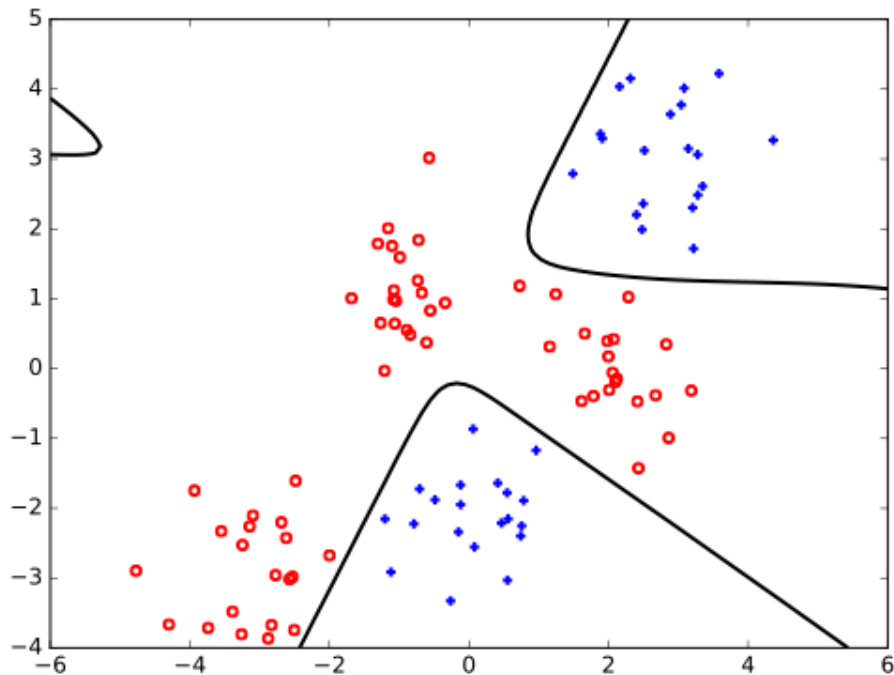


100 hidden units



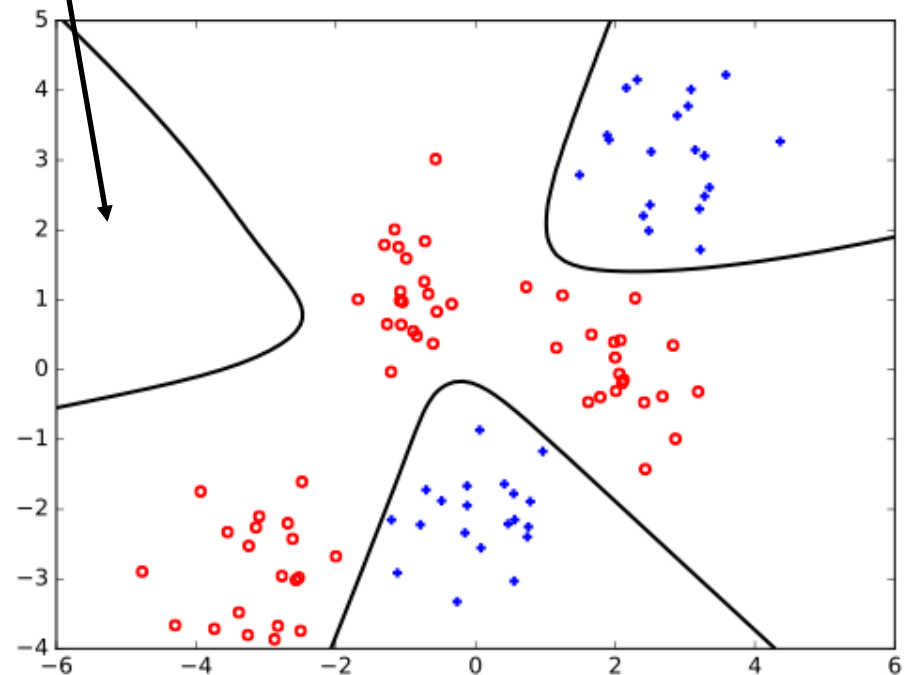
Decisions (and a harder task)

10 hidden units



???

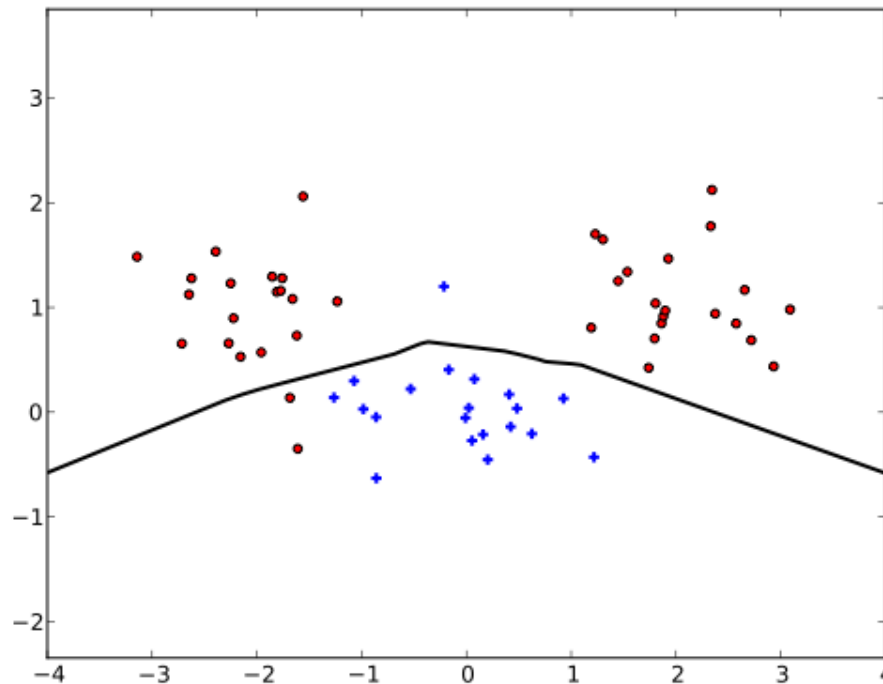
100 hidden units



Architectural Variations...

- Increasing the number of hidden units
 - More powerful decision boundary
 - Easier to fit the data

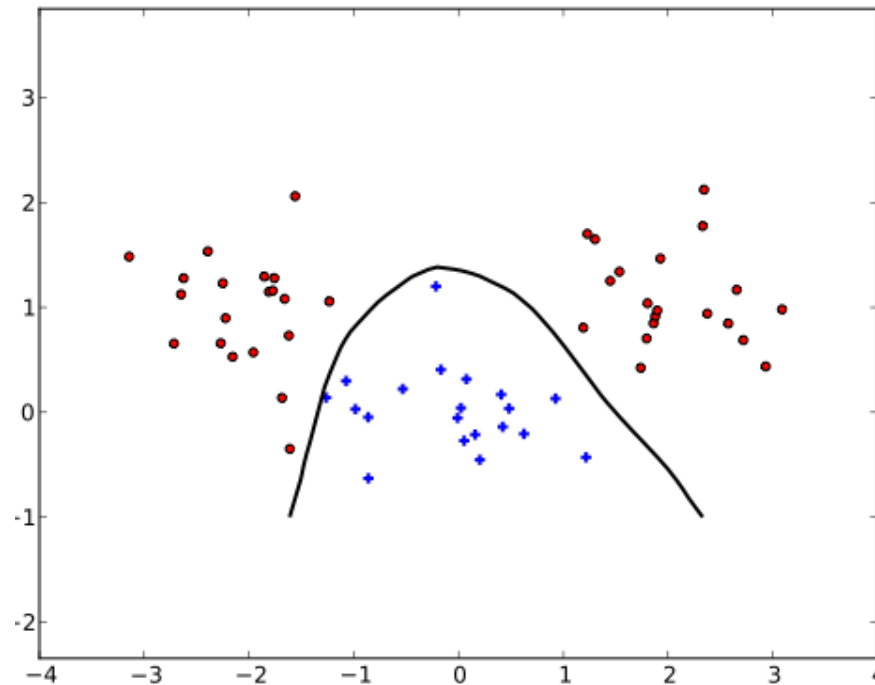
10 hidden units



Architectural Variations...

- Increasing the number of hidden units
 - More powerful decision boundary
 - Easier to fit the data

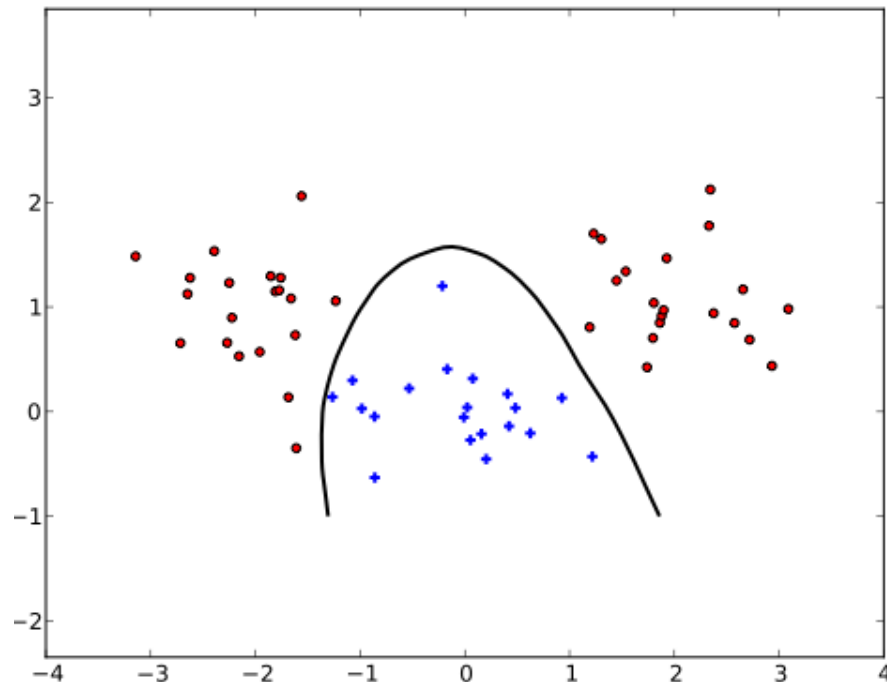
100 hidden units



Architectural Variations...

- Increasing the number of hidden units
 - More powerful decision boundary
 - Easier to fit the data

500 hidden units



Representational Power

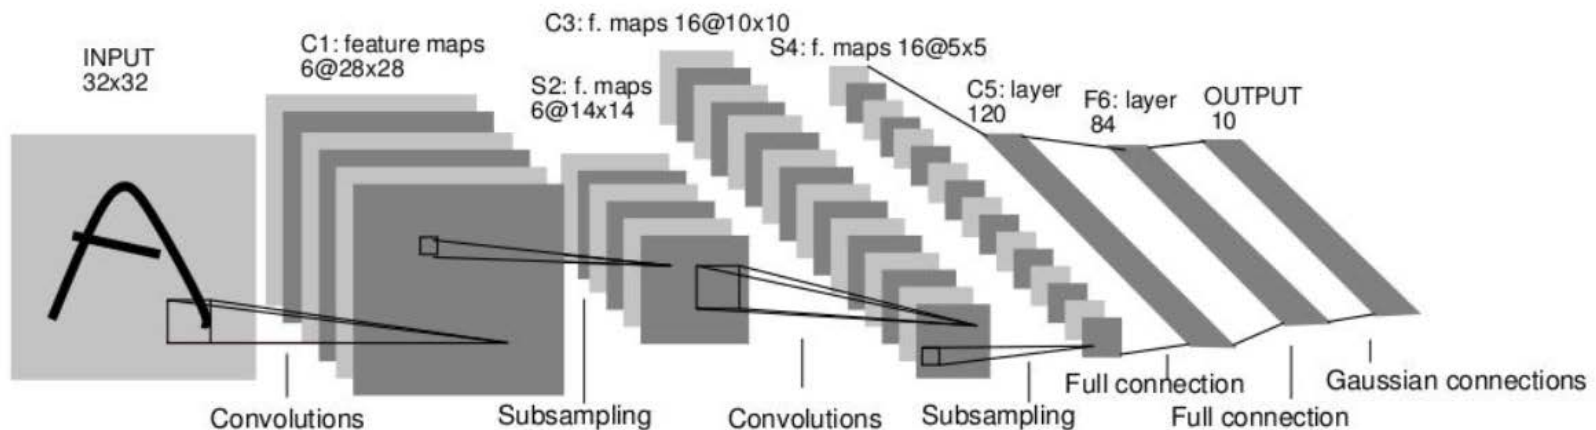
- 1 layer: Linear decision surface
- 2+ layers can represent any function assuming non-trivial non-linearity
- For fully connected models 2 or 3 layers seems the most that can be effectively trained
- Regarding number of units/layer:
 - Parameters grows with $(\text{units/layer})^2$

Architectures

- How to select:
 - Depth
 - Width
 - Parameter count
- Manual selection of features has turned into manual selection of architectures

Convolutional Neural Networks

- LeCun et al. 1989
- Neural network with specialized connectivity structure



Application to ImageNet

- ~14 million images, 20K classes
- Images gathered from Internet
- Human labeled via Amazon Mechanical Turk



Goal

- Image recognition
- Image \rightarrow class label



lens cap

reflex camera
Polaroid camera
pencil sharpener
switch
combination lock



abacus

abacus
typewriter keyboard
space bar
computer keyboard
accordion



slug

slug
zucchini
ground beetle
common newt
water snake

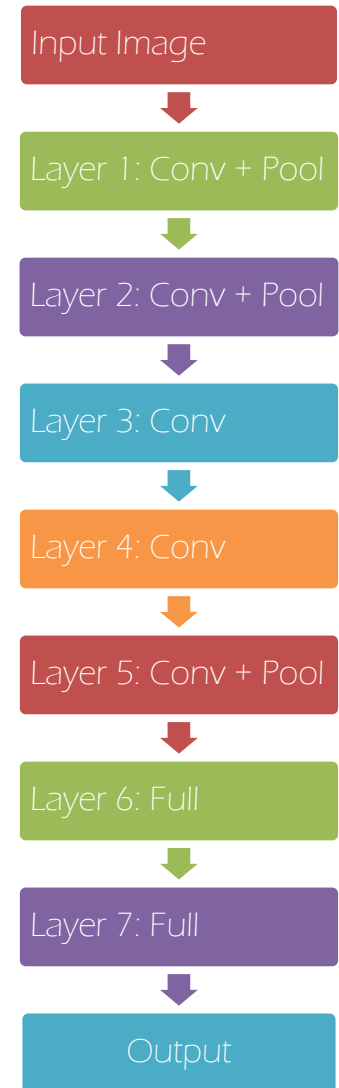


hen

hen
cock
cocker spaniel
partridge
English setter

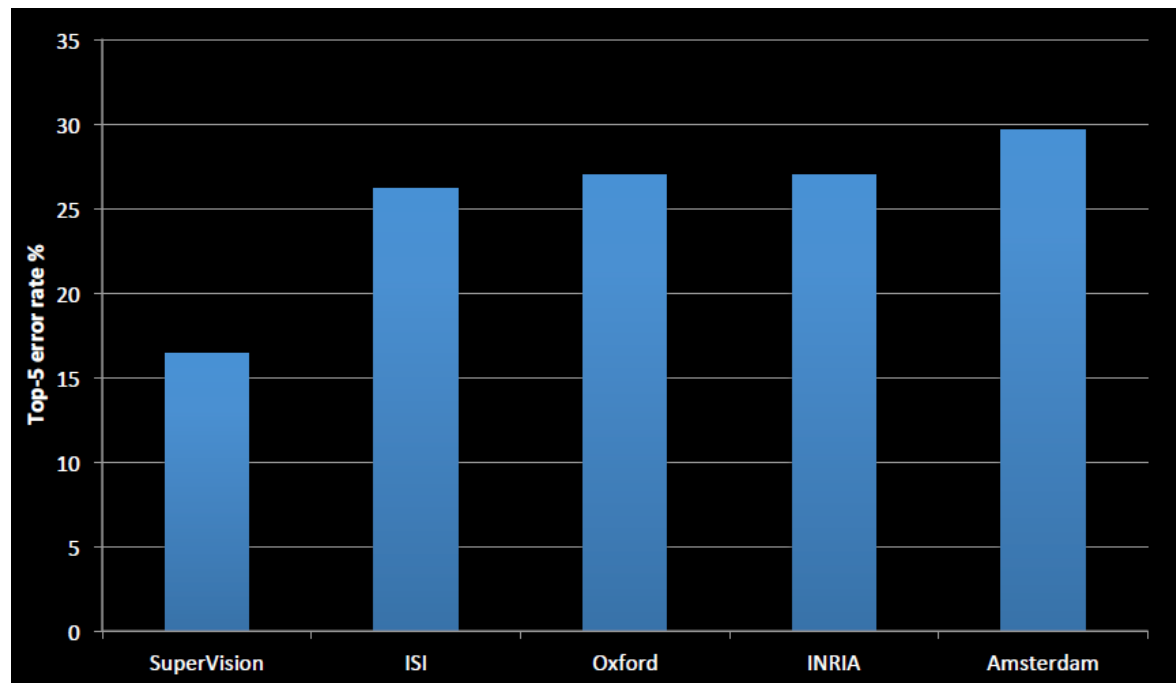
Krizhevsky et al. [NIPS 2012]

- Same model as LeCun'98 but
 - Bigger (8 layers)
 - 7 hidden layers, 650K neurons, 60M parameters
 - More data (10^6 vs 10^3 images)
 - GPU implementation (50x speedup over CPU)
 - Trained on 2 GPUs for a week
 - Better regularization



ImageNet Classification 2012

- Krizhevsky et al. – 16.4% error (top-5)
- Next best (non-convent) – 26.2% error

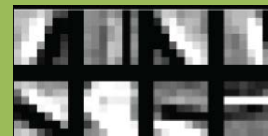


Components of Each Layer

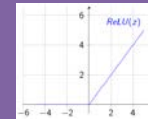
Pixels/Features



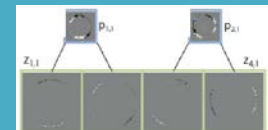
Filter with learned dictionary



Non-linearity



Spatial local max pooling



[Optional] Normalization across data/features



Output Features

Filtering

- Convolutional
 - Dependencies are local
 - Translation invariance
 - Tied filter weights (few parameters)



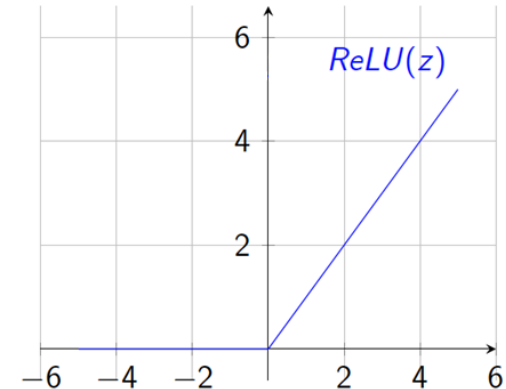
Input



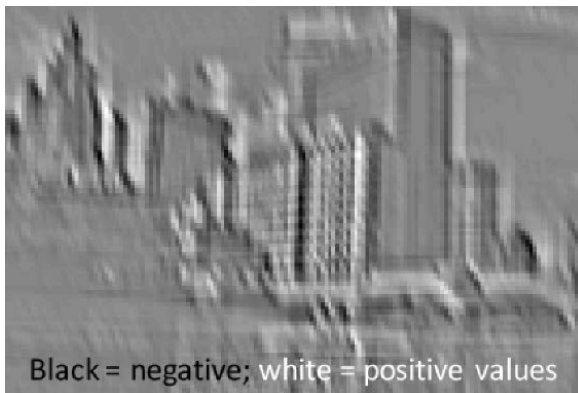
Feature Maps

Non-linearity

- Rectified linear function
 - Applied per-pixel
 - Output = $\max(0, \text{input})$



Input feature map

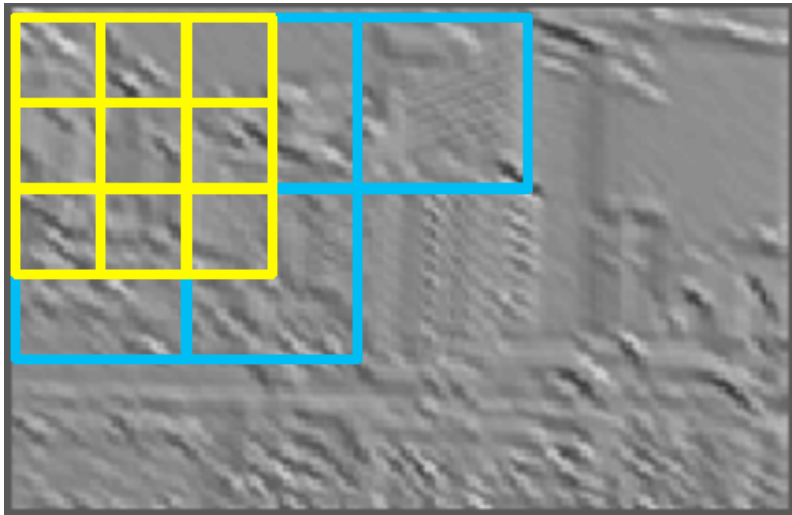


Output feature map

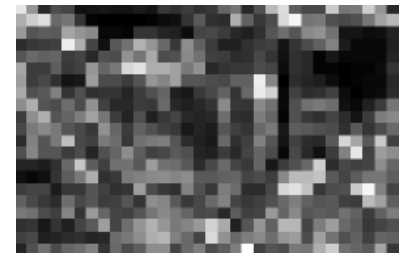


Pooling

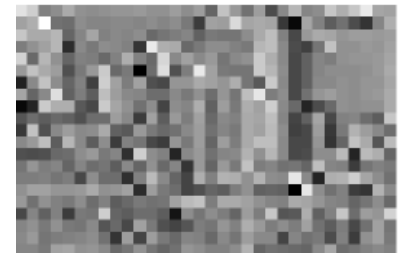
- Spatial Pooling
 - Non-overlapping/overlapping regions
 - Sum or max
 - Invariance to small transformations
 - Larger receptive field



Max

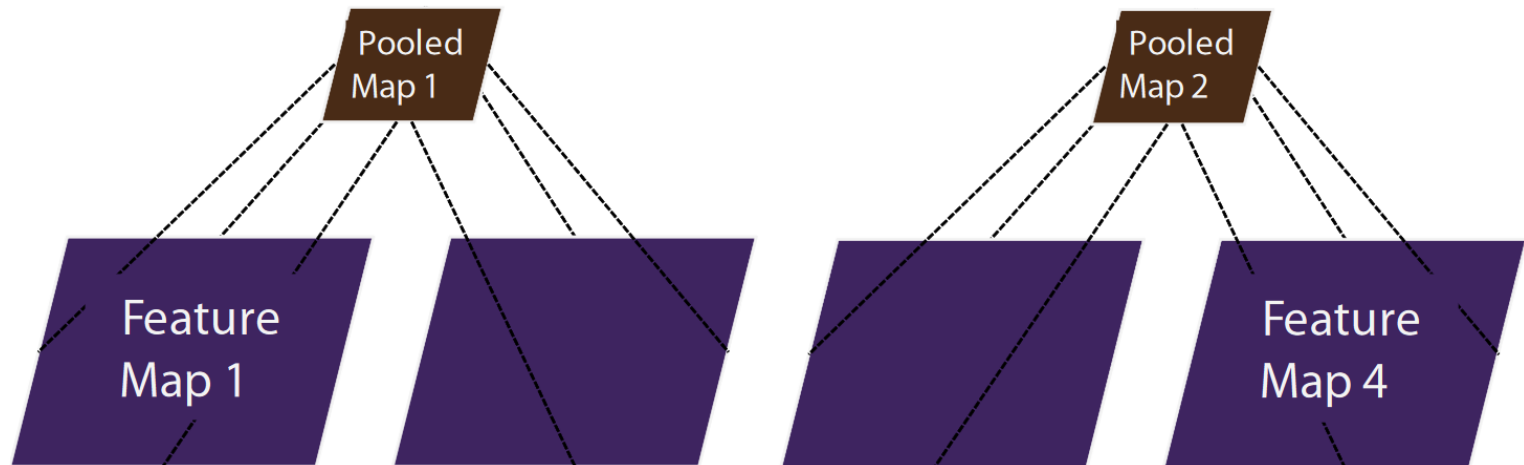


Sum



Pooling

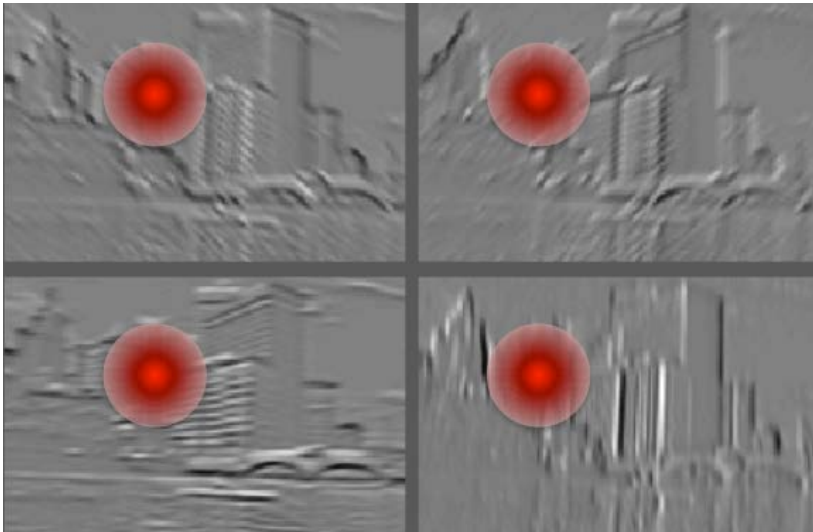
- Pooling across feature groups



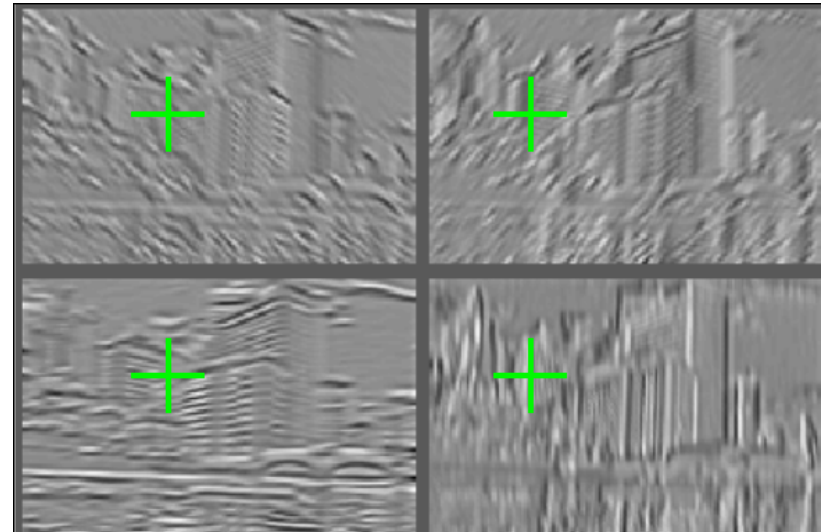
Normalization

- Contrast normalization (across feature maps)

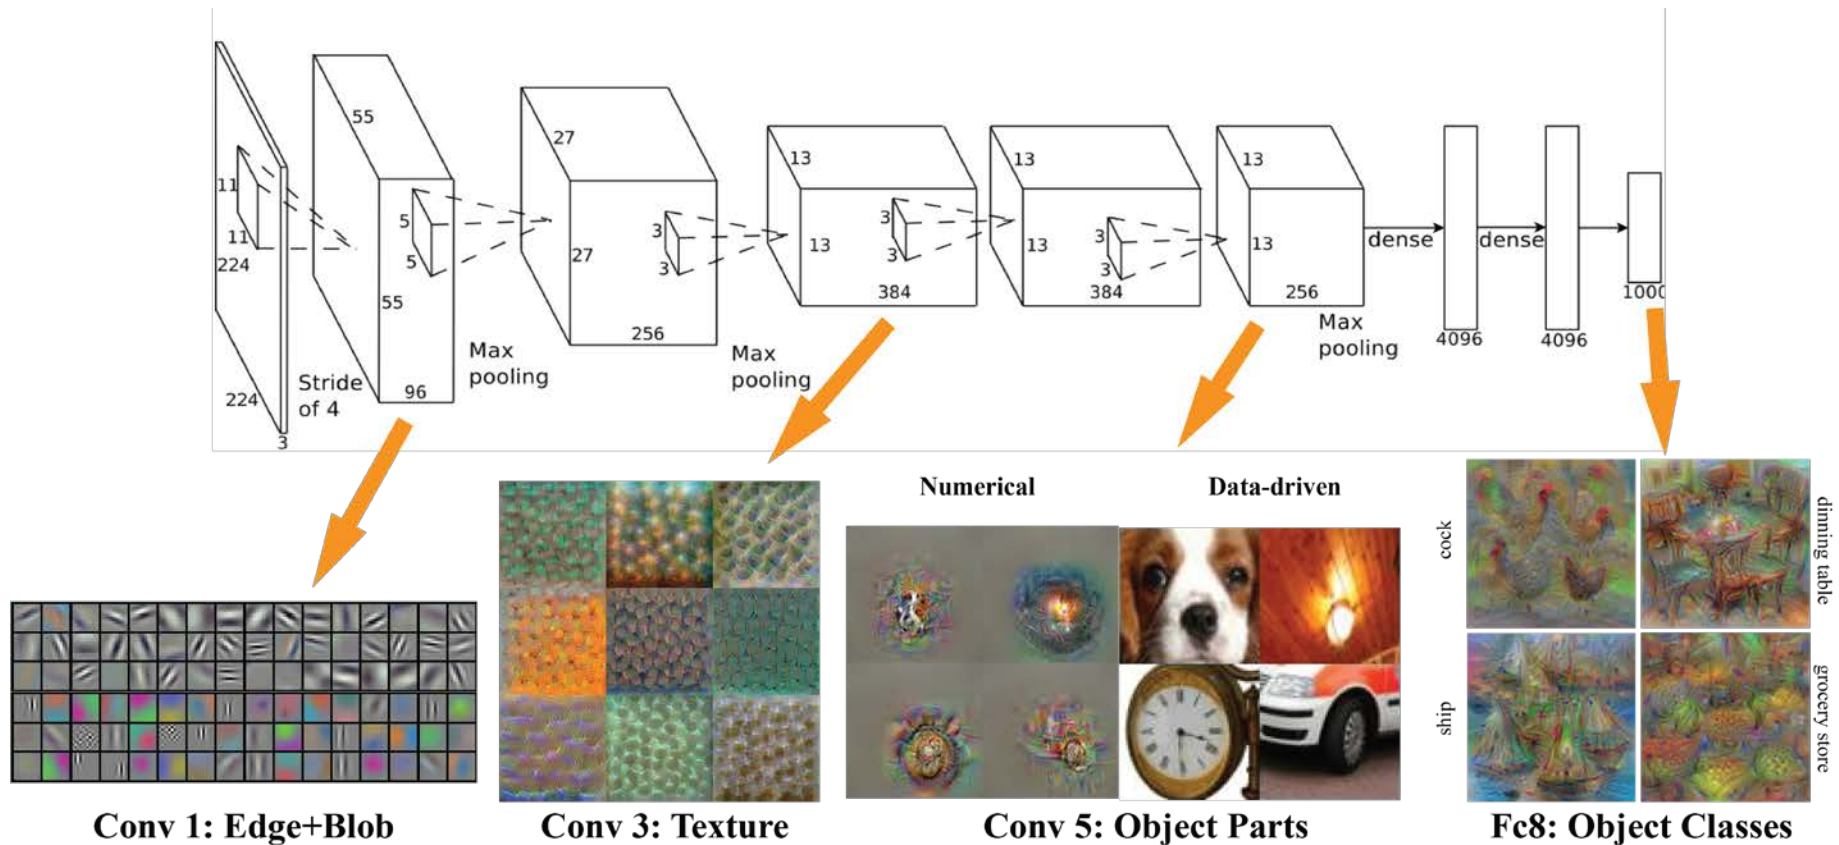
feature map



feature map after normalization



AlexNet Architecture



Summary

- Feedforward neural networks can be trained effectively using SGD
- Back-propagation
 - relies on applying chain rule to compute error gradient with respect to the network parameters
- CNN – neural network with a specialized connectivity structure
 - state-of-the-art for many computer vision problems