

6.036 Introduction to Machine Learning

(meets with 6.862)

**Generalization, Complexity,
VC-dimension (Chapter 7 in notes)**

Administrivia

Project #2 due Friday 4/7 @ 9AM.

As always:

- Check LMOD/Piazza for announcements.
- To contact staff, use Piazza
(6036-staff@lists.csail.mit.edu for exceptions only)

So far...

► **Tasks:**

- Binary or multi-class classification

$$h\left(\text{BEHIND EVERY CODE IS AN ENIGMA THE IMITATION GAME}$$

- Regression

$$h\left(\text{LIVING ROOM WITH LARGE WINDOWS AND SOFA}$$

- **Formulations:** hyperplane separation, empirical risk minimization, regularization, neural networks,...
- **Algorithms:** perceptron, stochastic gradient, backprop,...

Generalization

Key notion in Machine Learning

- Find classifier that does well on **training set**
- Expect (hope?) that also does well on **test set** (validation set, population)

The real goal is to perform well on the **test set** (population), not just on the **training set**

How to ensure generalization?

Today, we'll study:

- How to formalize learning/generalization
- Role of parameters (what controls generalization)?
- How to measure the size of a model class

Our setting


Assume a probabilistic model, where **both** training and validation sets come from the **same** distribution \mathbf{P}^*

- **Example:** mixture of Gaussians

$$X \sim p_1 \mathcal{N}(\mu_1, \Sigma_1) + p_2 \mathcal{N}(\mu_2, \Sigma_2)$$





- From this:



Training set

Validation set



Measuring errors

- Distribution \mathbf{P}^* is fixed, but unknown
- Generates both the training and validation sets

How to measure the errors of a given classifier $h(x)$?

- On the training set, **empirical error** (**# of mistakes**)

$$\mathcal{E}_n(h) = \frac{1}{n} \sum_{t=1}^n [[y^{(t)} h(x^{(t)}) \leq 0]]$$

Depends on
examples!

- On the population, **test** or **generalization error**

$$\mathcal{E}(h) = \mathbf{E}_{(x,y) \sim \mathbf{P}^*} [[yh(x) \leq 0]]$$

Does not,
population only

(notice, this is also the prob. of h misclassifying a random point)

Empirical risk minimization

- Ideally, would like to directly minimize test error $\mathcal{E}(h)$

$$h^* = \arg \min_{h \in \mathcal{H}} \mathcal{E}(h)$$

but we don't have direct access to this... :(

- Instead, minimize training (empirical) error

$$\hat{h} = \arg \min_{h \in \mathcal{H}} \mathcal{E}_n(h)$$


Q: When does small $\mathcal{E}_n(h)$ imply small $\mathcal{E}(h)$?

Q: What is the role of the model class \mathcal{H} ?

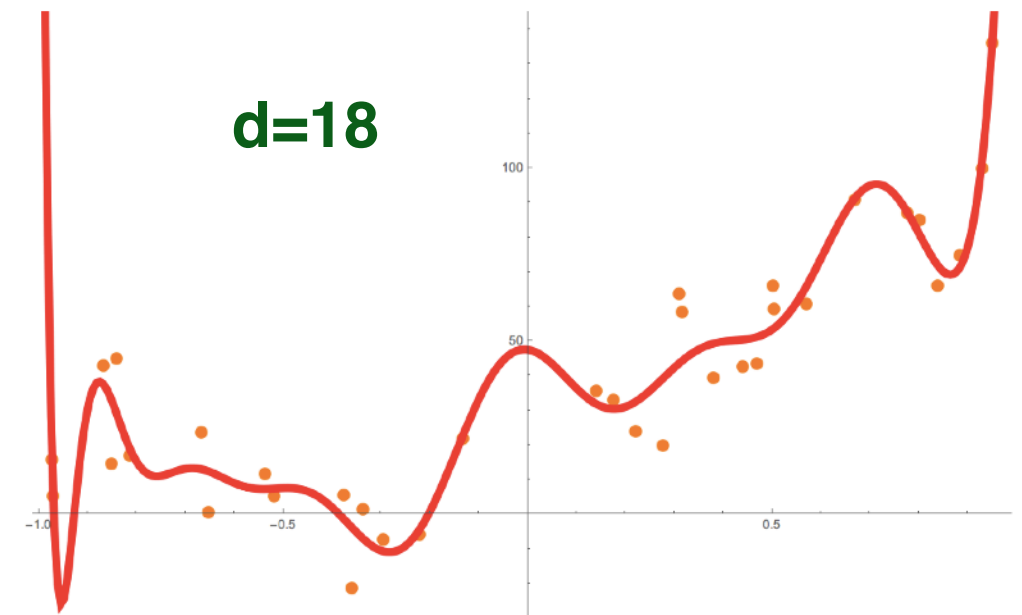
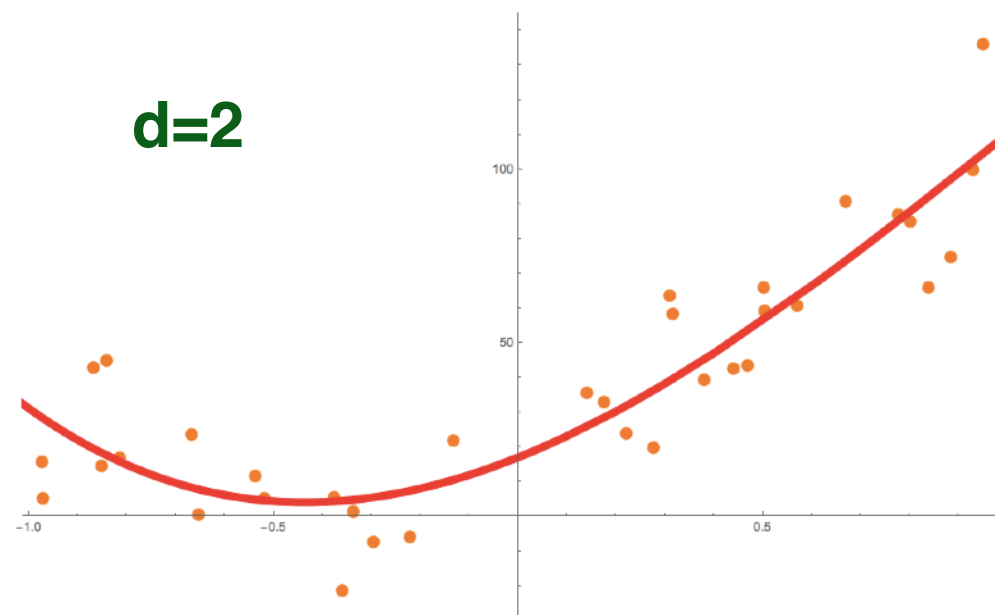
Model class and generalization

- ▶ **Intuition:** “small” model classes avoid fitting noise

- ▶ From classification lecture:



- ▶ From regression lecture:



Generalization guarantees

- What kind of guarantees can we expect?
 - Want small test/generalization error $\mathcal{E}(h)$
 - But, recall that training set is random (could be very bad — though hopefully with small probability)
- Probably Approximately Correct (PAC) Framework

With probability at least $1-\delta$ over the training set (sampled from \mathbf{P}^), the generalization error satisfies:*

$$\mathcal{E}(\hat{h}) \leq \epsilon$$

Generalization

- ▶ What must the relationship between all these parameters be, if we want to ensure learning?

$$\delta, \quad \epsilon, \quad n, \quad |\mathcal{H}|$$

confidence, error bound, number of samples, size of model class

- ▶ Assumptions (for simplicity only, can be removed):
 - Finite set of classifiers ($|\mathcal{H}|$ finite)
 - There exists at least one perfect classifier h^* (*realizable* case, i.e., $\mathcal{E}(h^*)=0$).

Simplified analysis

- ▶ Recall PAC: *With probability at least $1-\delta$ over the training set (sampled from \mathbf{P}^*), the generalization error is below ϵ .*
- ▶ Fix ϵ , \mathcal{H} , and n .
How small can failure probability δ be?
(to avoid ERM going badly on training set)
- ▶ Pick a “bad” classifier (i.e., one with $\mathcal{E}(h) > \epsilon$)
 - Could ERM conceivably choose this as a solution (i.e., could h behave well on training set)?
 - Certainly, if h correctly classifies *every* sample: $(1 - \epsilon)^n$
- ▶ Since there may be many such classifiers:

$$\delta \leq |\mathcal{H}| (1 - \epsilon)^n$$

Generalization bound

$$\delta \leq |\mathcal{H}| (1 - \epsilon)^n$$

- ▶ From this, simple manipulations yield:

$$\log \delta \leq \log |\mathcal{H}| + n \log(1 - \epsilon) \leq \log |\mathcal{H}| - n\epsilon$$

- ▶ and rearranging

$$\epsilon \leq \frac{\log |\mathcal{H}| + \log(1/\delta)}{n}$$

- ▶ Insights:

- For good generalization, $|\mathcal{H}|$ must be “small”
- The more confident we want to be (smaller δ), the more training samples we’ll need.
- Increasing number of samples n decreases error

More generally...


- This analysis is somewhat restricted (“realizable case”).
- Can be extended to the general situation:

$$\mathcal{E}(\hat{h}) \leq \mathcal{E}_n(\hat{h}) + \sqrt{\frac{\log |\mathcal{H}| + \log(2/\delta)}{2n}}$$

- Similar qualitative dependence on parameters

Infinite model classes

- In some situations, $|\mathcal{H}|$ is infinite!
 - E.g., set of all linear classifiers



- How to quantify the *size* of the model class?
(naive “count” gives infinity!)
- Insight: only thing that matters are the *labelings* induced by a given classifier

Growth function

- Consider the labels induced by all classifiers in \mathcal{H}

$$\begin{array}{ccccc} & x^{(1)} & x^{(2)} & \dots & x^{(n)} \\ h \in \mathcal{H} : & +1 & -1 & \dots & -1 \\ h' \in \mathcal{H} : & +1 & -1 & \dots & -1 \\ h'' \in \mathcal{H} : & +1 & +1 & \dots & -1 \\ \dots & \dots & \dots & \dots & \dots \end{array}$$

and let $N_{\mathcal{H}}(x^{(1)}, \dots, x^{(n)})$ be the number of distinct rows

- This depends on the examples, so maximize

$$N_{\mathcal{H}}(n) = \max_{x^{(1)}, \dots, x^{(n)}} N_{\mathcal{H}}(x^{(1)}, \dots, x^{(n)})$$

- This is the *growth function*, and quantifies how powerful our family \mathcal{H} of classifiers is

Generalization guarantees

- › Similarly as before, we can give guarantees, but now in terms of the growth function:

$$\mathcal{E}(\hat{h}) \leq \mathcal{E}_n(\hat{h}) + \sqrt{\frac{\log N_{\mathcal{H}}(2n) + \log(4/\delta)}{n}}, \text{ for all } \hat{h} \in \mathcal{H}$$

- › When does this go to zero as n increases?
- › Need $\log N_{\mathcal{H}}(2n)$ to grow slower than n

VC-dimension

- A measure of complexity of a set of classifiers, the Vapnik-Chervonenkis (VC) dimension d_H
 - Largest number of points that can be labeled in all possible ways using classifiers from \mathcal{H}
 - Equivalently, largest n for which $N_{\mathcal{H}}(n) = 2^n$
- If $n > d_H$, then the number of labelings grows only logarithmically.


$$\log N_{\mathcal{H}}(2n) \leq d_{\mathcal{H}}(\log(2n/d_{\mathcal{H}}) + 1)$$

- Even if set of classifiers is infinite, finite VC-dimension can guarantee generalization


(Other existing approaches: regularization, Rademacher complexity, etc...)

VC-dimension (example)

- ▶ **Example:** in \mathbb{R}^2 , can arbitrarily label 3 points



but not 4 points:



- ▶ For *linear classifiers* in \mathbb{R}^d , the VC-dimension is *exactly* $d+1$.

Summary - Generalization

- For learning, need to control **generalization error**, not just training error.
- PAC (probably approximately correct) framework:
with high probability, a classifier that has small training error will have small generalization error.
- For generalization, it is sufficient to restrict the size of the model class \mathcal{H}
- For uncountable model classes, can quantify size using VC-dimension