

EM (cont'd) and Hidden Markov Models



Latent variable models

- Latent variable models allow us to hypothesize different types of structures that may underlie the observed data and then recover those structures
- We have understand how to
 - specify them (variables, distributions with parameters)
 - sample from them (as generative models)
 - estimate them from data
- Recall mixture models

$$\begin{array}{lll} y & y \in \{1, \dots, K\} & P(y) = p_y \\ \downarrow & & \\ x & x \in \mathbb{R}^d & P(x|y) = N(x; \mu^{(y)}, \sigma_y^2 I) \end{array}$$

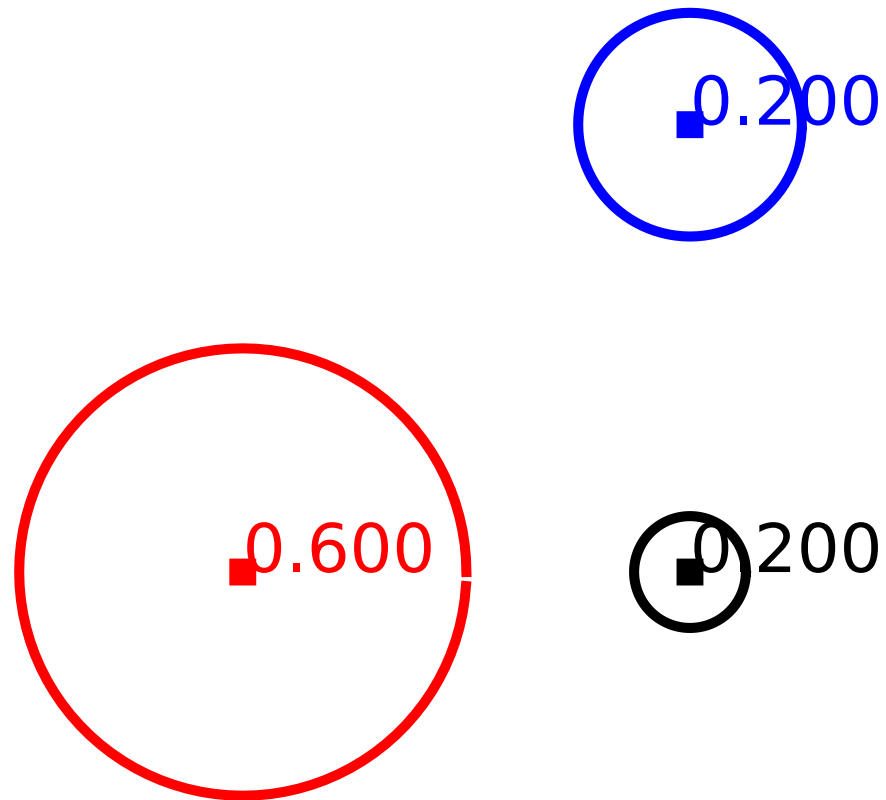


Mixture model: generation

► We can sample points from a mixture model in two steps

1) Sample $y \sim \text{Categ}(p_1, \dots, p_K)$ **which cluster**

2) Sample $x \sim N(x; \mu^{(y)}, \sigma_y^2 I)$ **which point from the chosen cluster**



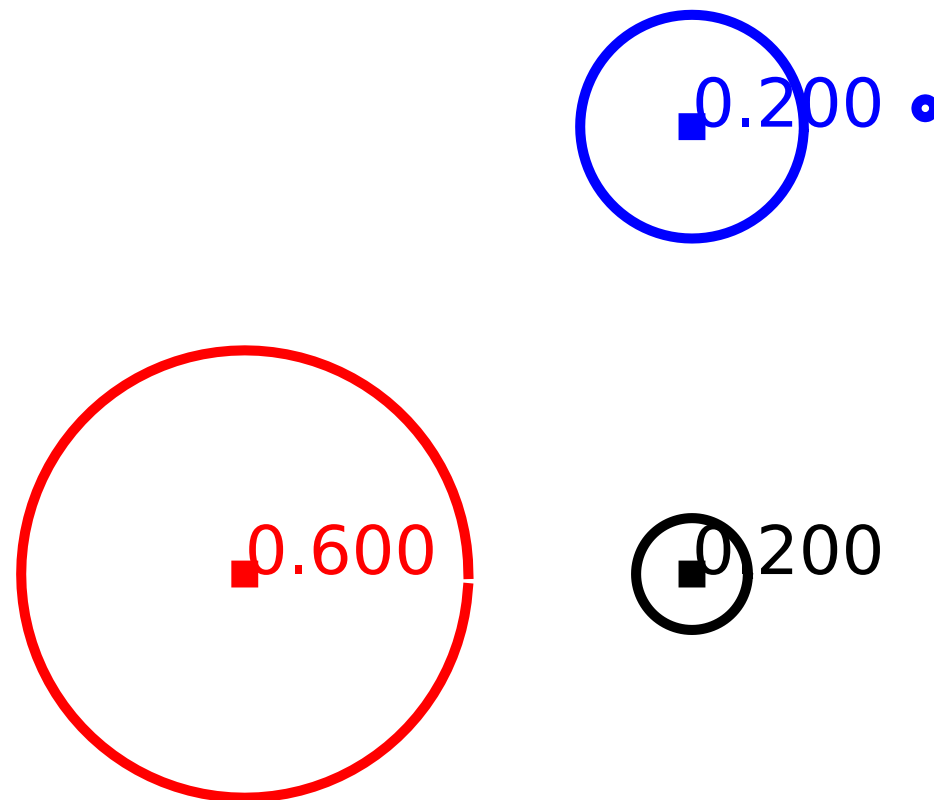


Mixture model: generation

► We can sample points from a mixture model in two steps

1) Sample $y \sim \text{Categ}(p_1, \dots, p_K)$ **which cluster**

2) Sample $x \sim N(x; \mu^{(y)}, \sigma_y^2 I)$ **which point from the chosen cluster**



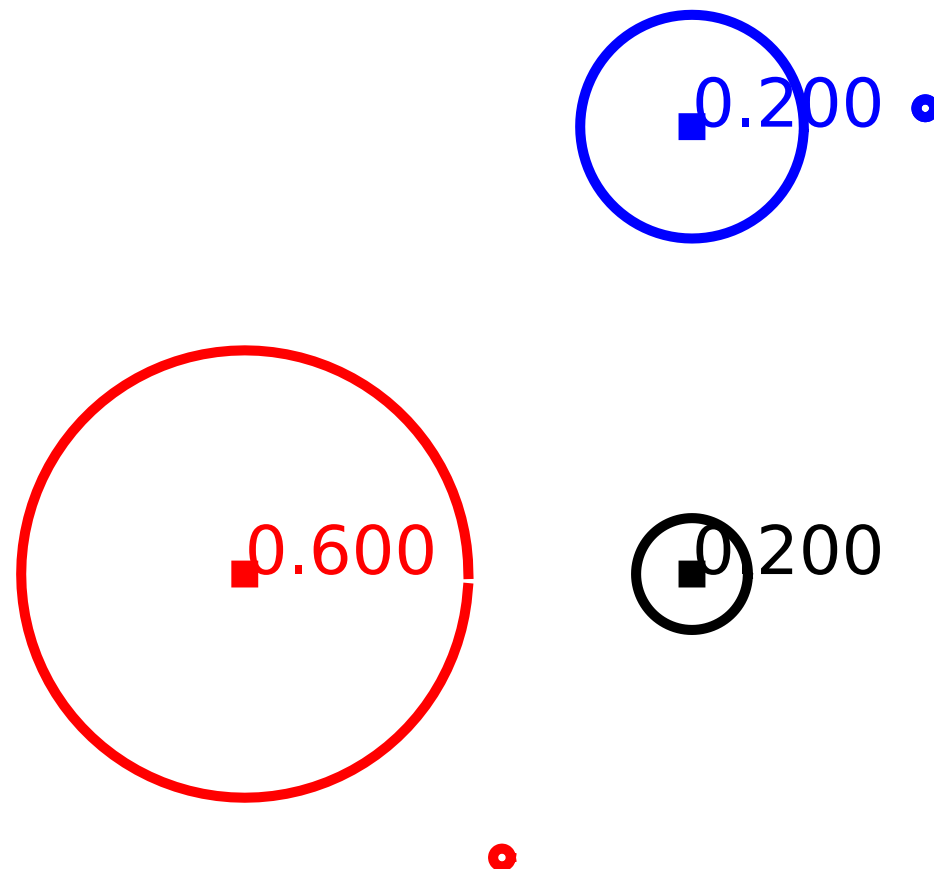


Mixture model: generation

► We can sample points from a mixture model in two steps

1) Sample $y \sim \text{Categ}(p_1, \dots, p_K)$ **which cluster**

2) Sample $x \sim N(x; \mu^{(y)}, \sigma_y^2 I)$ **which point from the chosen cluster**



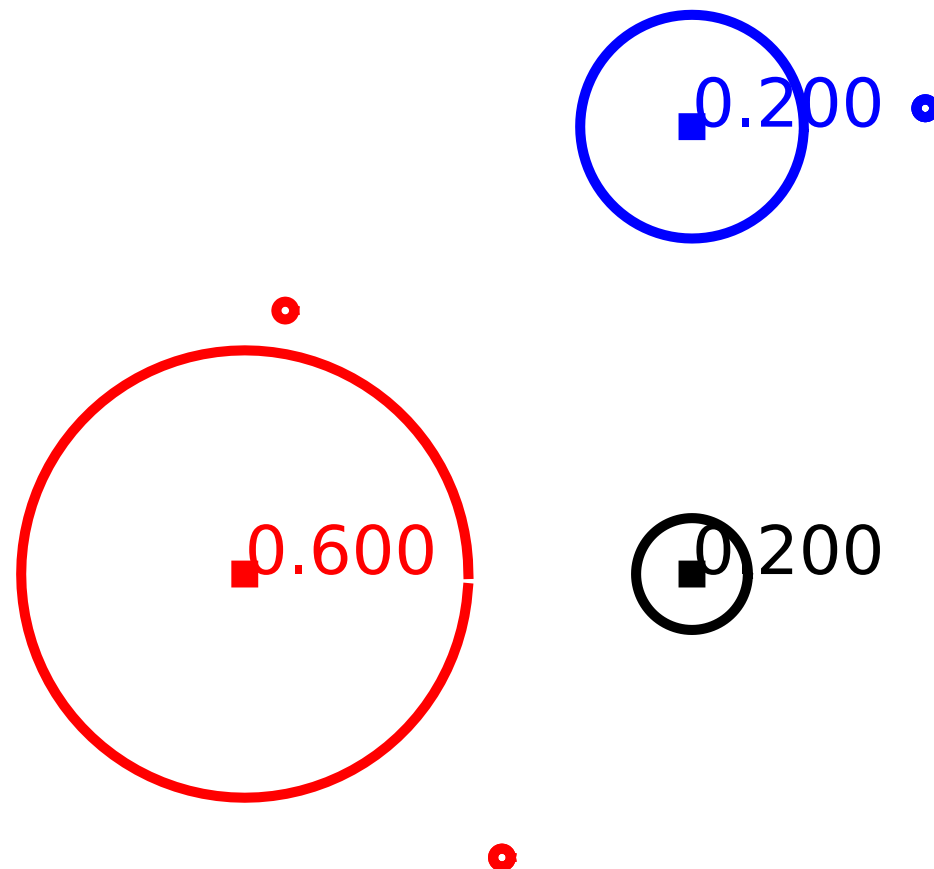


Mixture model: generation

► We can sample points from a mixture model in two steps

1) Sample $y \sim \text{Categ}(p_1, \dots, p_K)$ **which cluster**

2) Sample $x \sim N(x; \mu^{(y)}, \sigma_y^2 I)$ **which point from the chosen cluster**

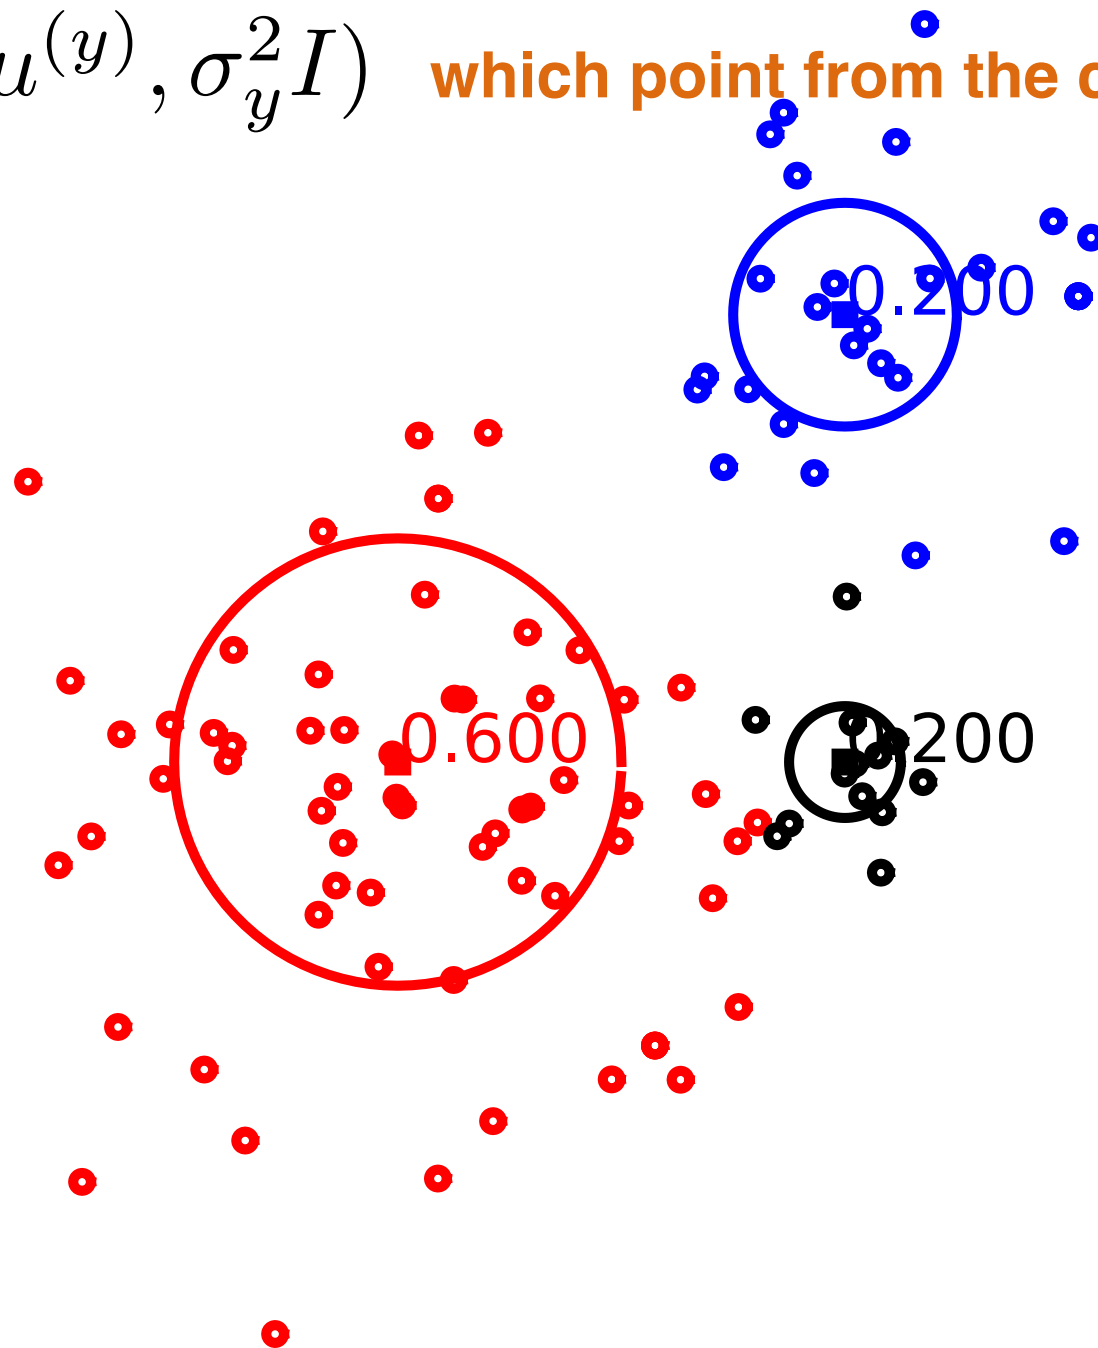


Mixture model: generation

► We can sample points from a mixture model in two steps

1) Sample $y \sim \text{Categ}(p_1, \dots, p_K)$ **which cluster**

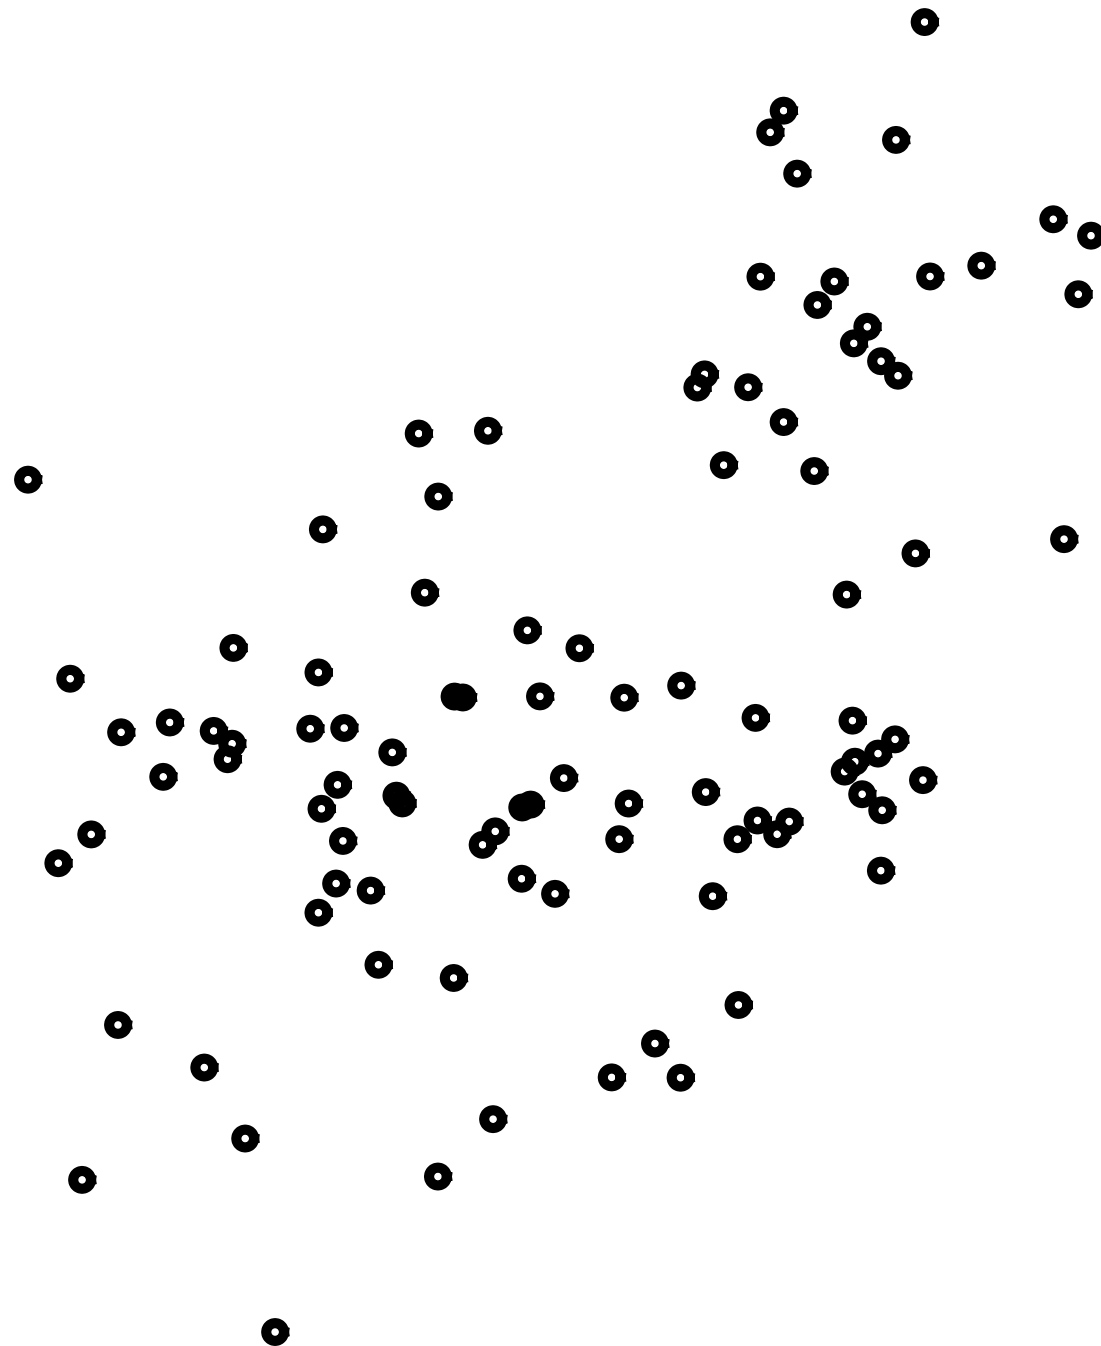
2) Sample $x \sim N(x; \mu^{(y)}, \sigma_y^2 I)$ **which point from the chosen cluster**





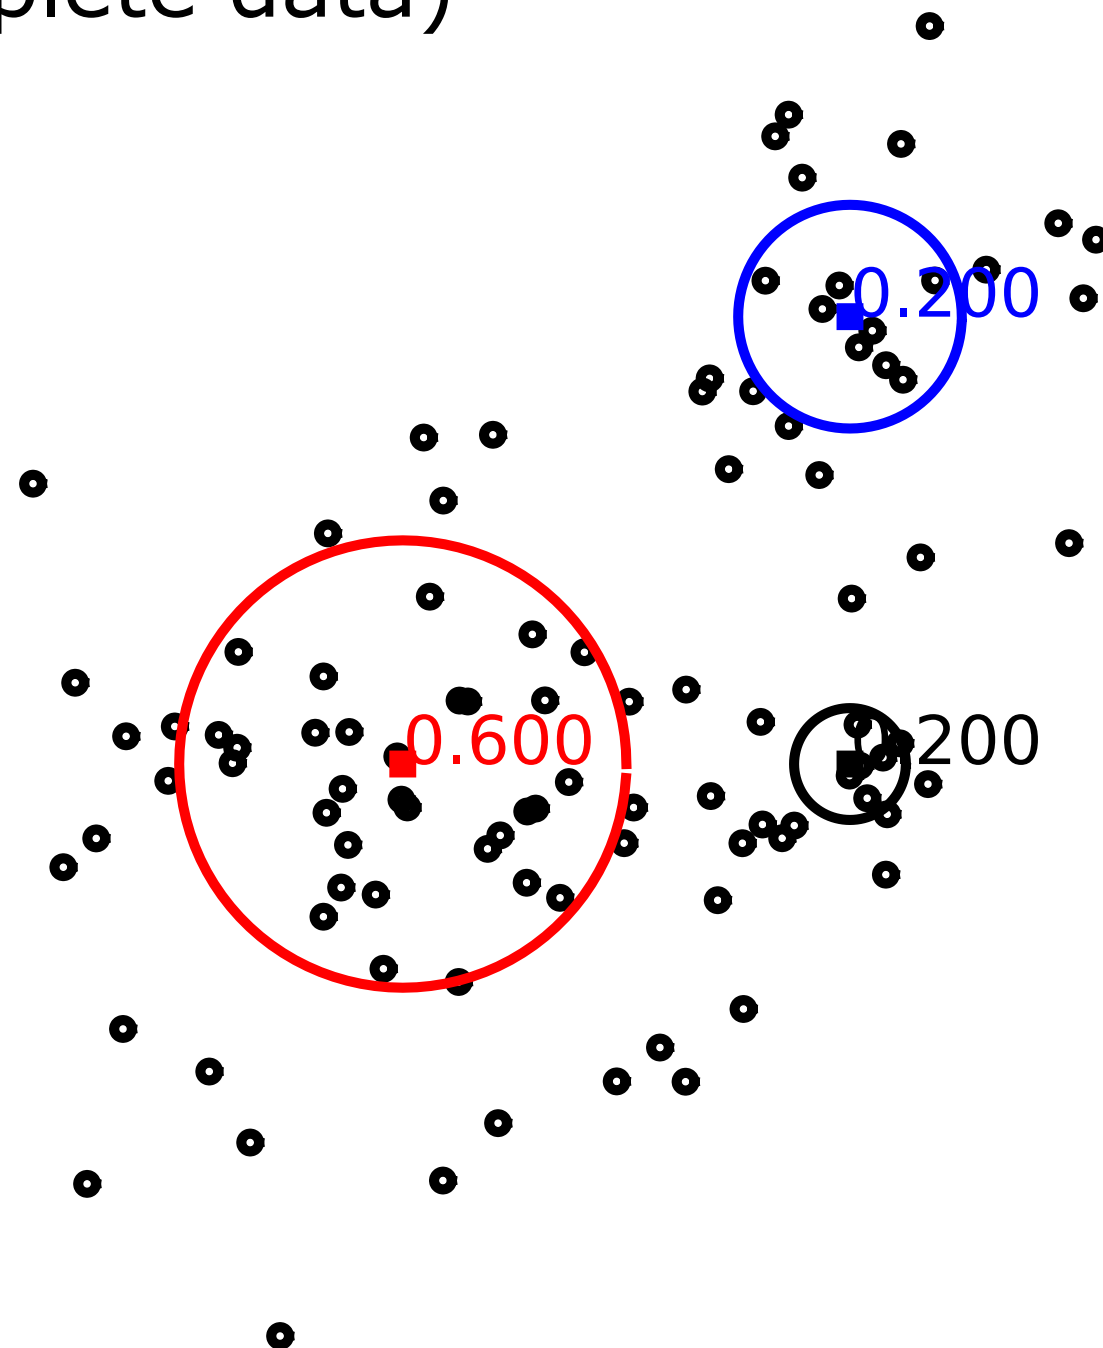
Mixture model: estimation

- But we typically don't get the cluster labels in the observed data, only the resulting points x



Mixture model: estimation

- Our goal is to try to estimate the underlying model (mixing proportions, component Gaussians) from the points alone (incomplete data)



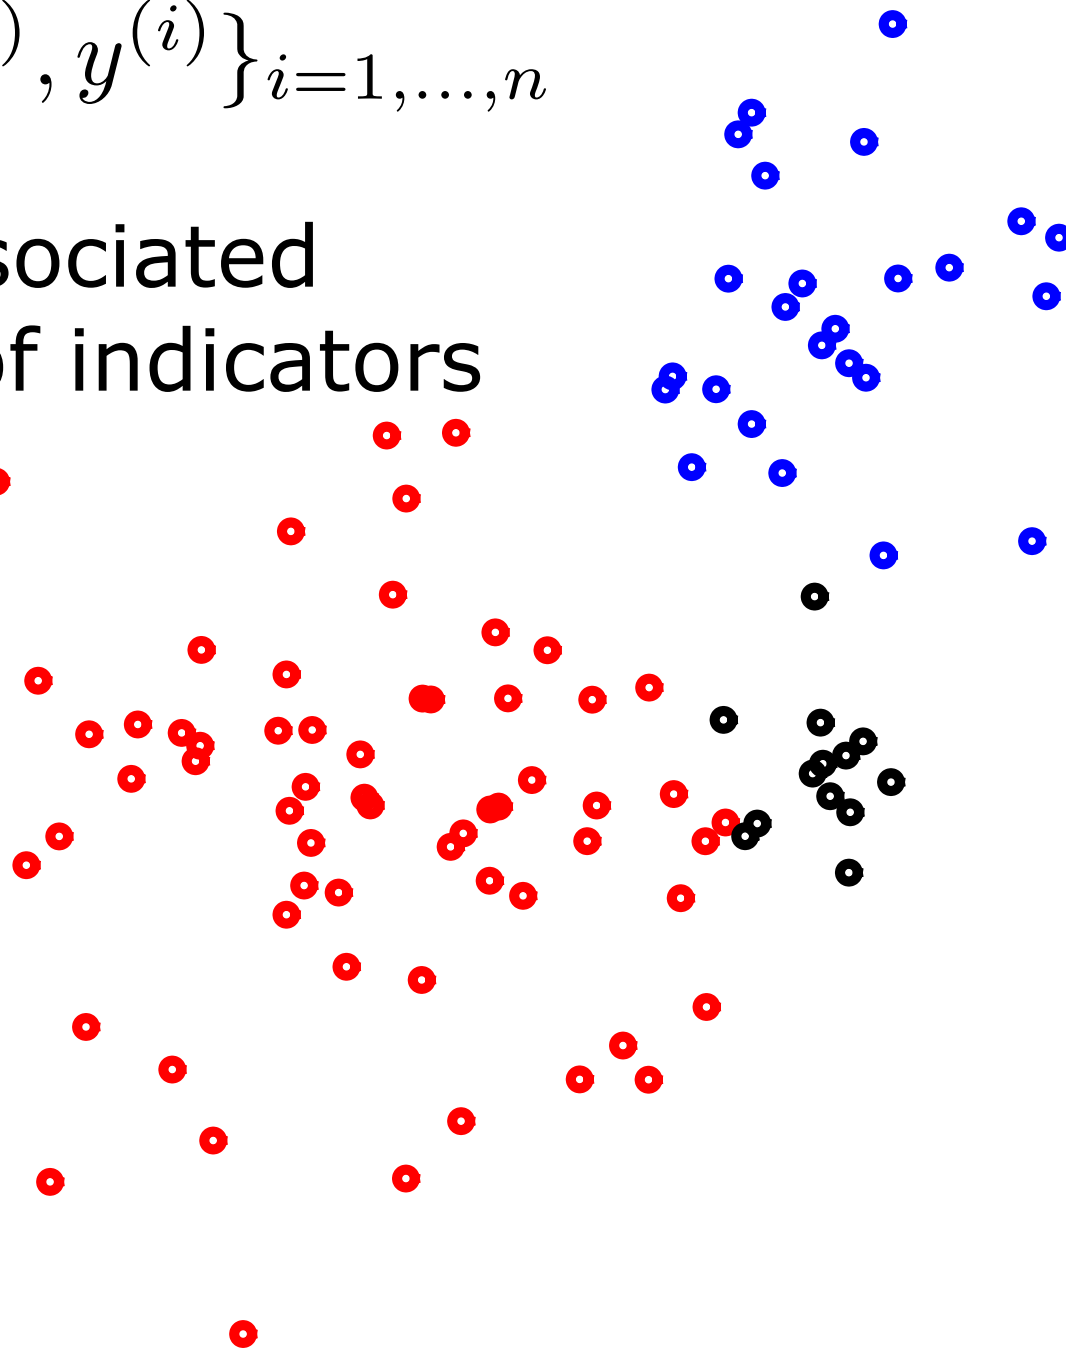
Complete data

- Estimation would be easy if we had complete data, i.e., pairs of cluster ids y and corresponding points x

$$\text{complete data} = \{x^{(i)}, y^{(i)}\}_{i=1, \dots, n}$$

- We represent the associated cluster ids in terms of indicators

$$\delta(y|i) = 1 \text{ if } y = y^{(i)} \\ \text{and } 0 \text{ otherwise}$$





Complete data: mixing proportions

- ▶ We maximize the log-likelihood of complete data

$$\sum_{y=1}^K \sum_{i=1}^n \delta(y|i) \log p_y$$

- ▶ Resulting ML estimate

$$\hat{p}_y = \frac{\sum_{i=1}^n \delta(y|i)}{n}, \quad y = 1, \dots, K$$



Complete data: Gaussians

- ▶ We maximize the log-likelihood of complete data (separately for each Gaussian)

$$\sum_{i=1}^n \delta(y|i) \log N(x^{(i)}; \mu^{(y)}, \sigma_y^2 I)$$

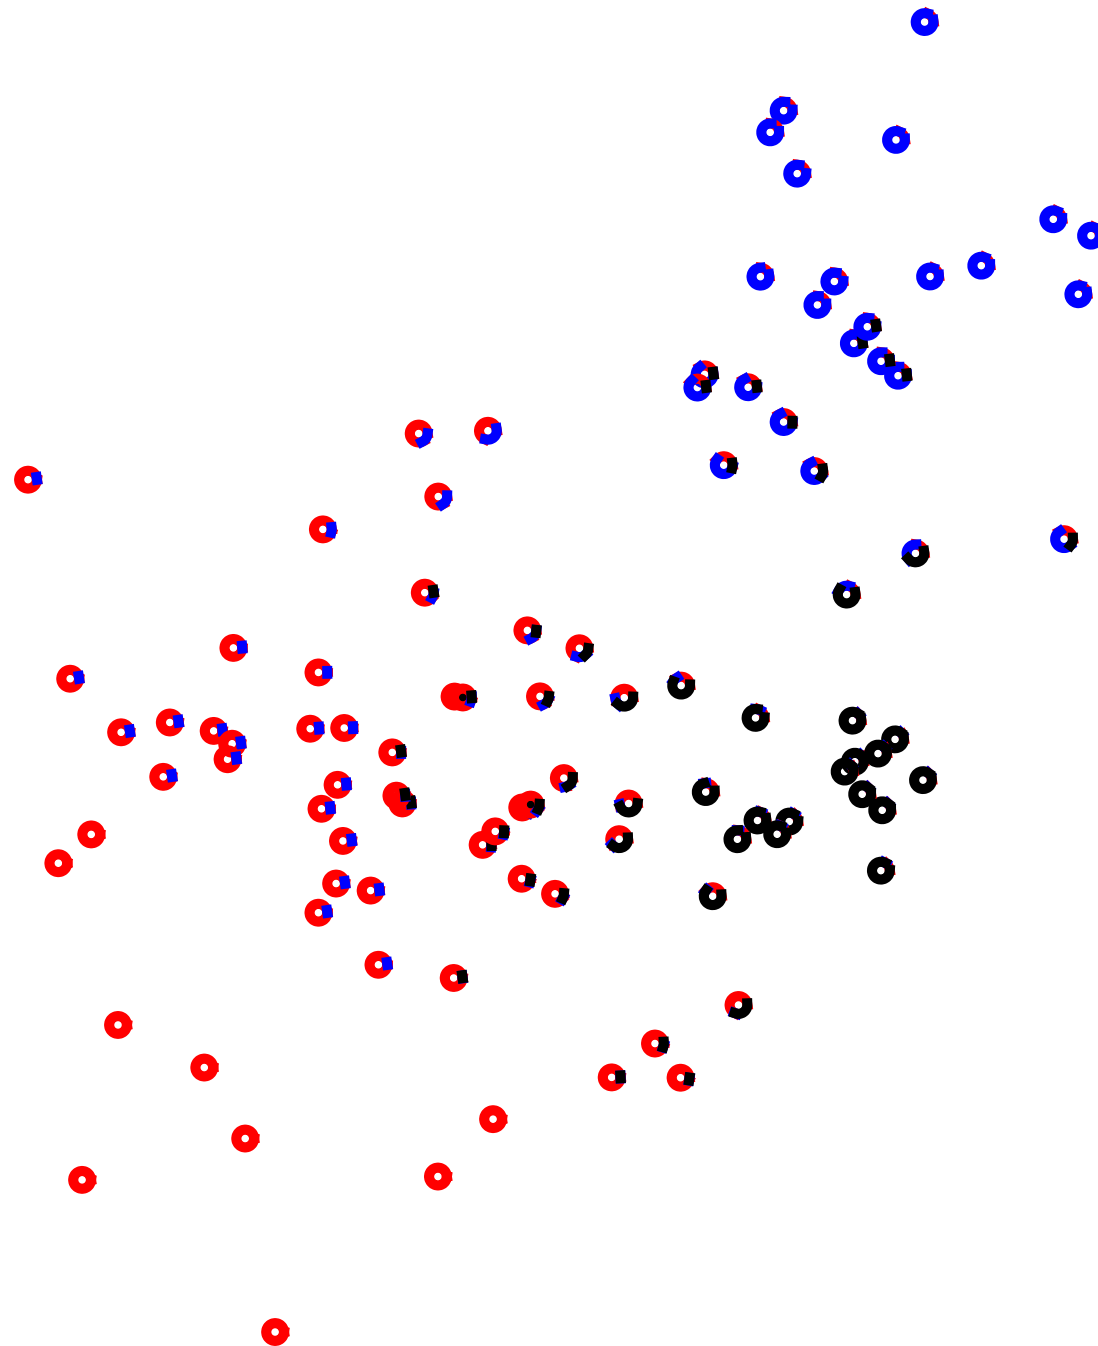
- ▶ The resulting ML estimates

$$\hat{\mu}^{(y)} = \frac{\sum_{i=1}^n \delta(y|i) x^{(i)}}{\sum_{i=1}^n \delta(y|i)}$$

$$\hat{\sigma}_y^2 = \frac{\sum_{i=1}^n \delta(y|i) \|x^{(i)} - \hat{\mu}^{(y)}\|^2}{d \sum_{i=1}^n \delta(y|i)}$$

Weighted data

- The estimation would be just as easy if instead of true cluster labels y , we had soft assignments



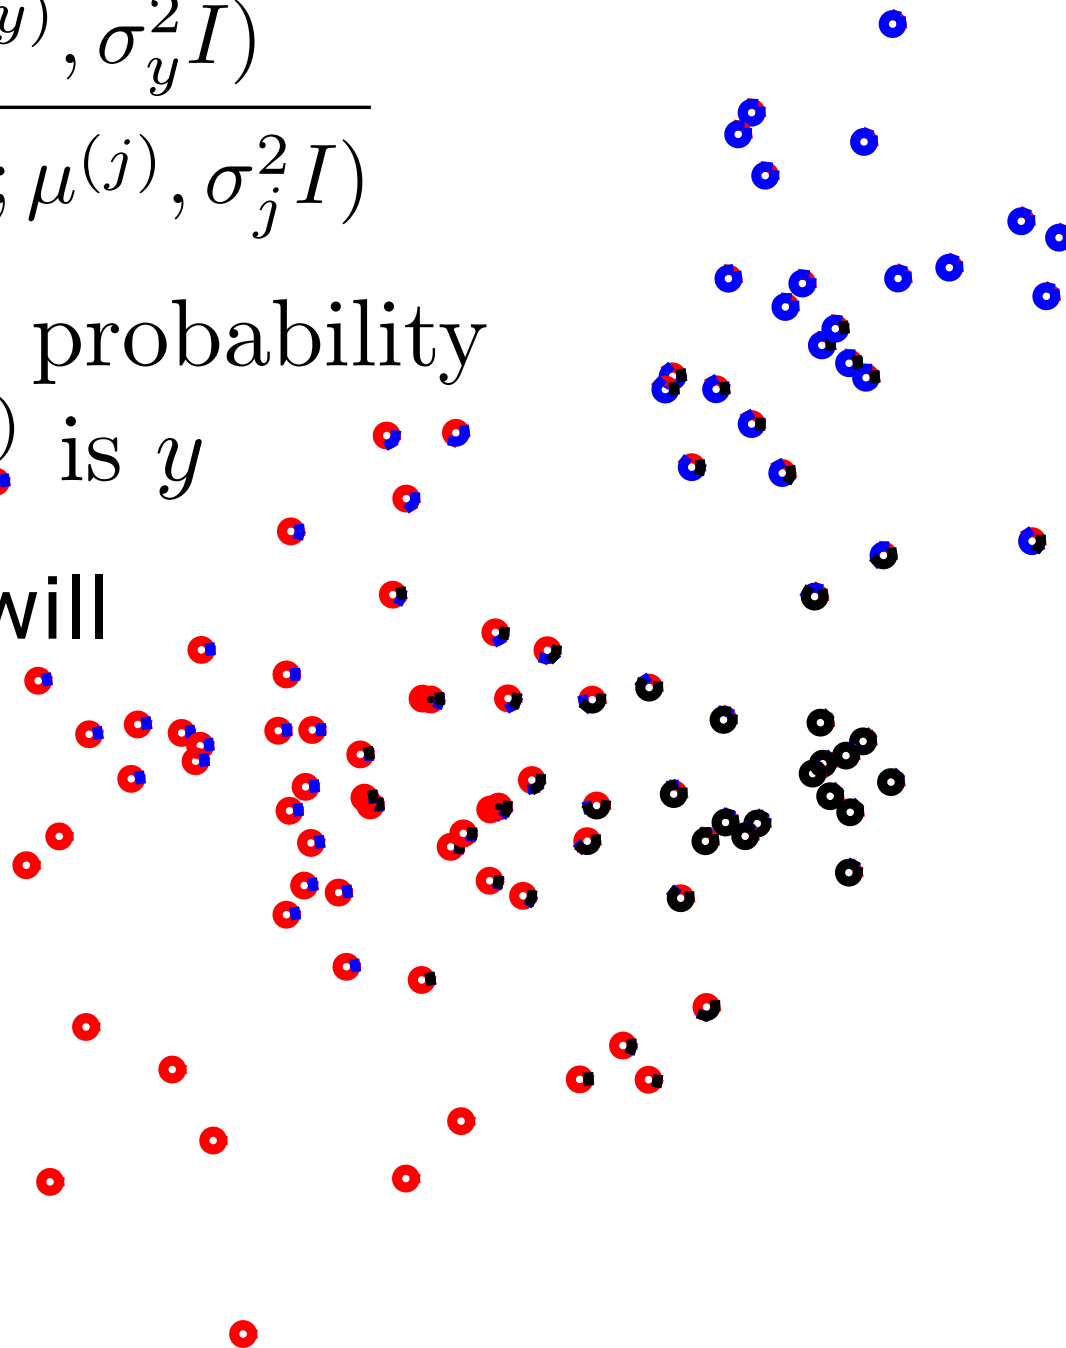
E-step: weighted data

- The estimation would be just as easy if instead of true cluster labels y , we had soft assignments

$$p(y|i) = \frac{p_y N(x^{(i)}; \mu^{(y)}, \sigma_y^2 I)}{\sum_{j=1}^k p_j N(x^{(i)}; \mu^{(j)}, \sigma_j^2 I)}$$

$p(y|i)$ is the posterior probability that the unknown $y^{(i)}$ is y

(these assignments will change as the model is updated)





M-step: mixing proportions

- We maximize the expected (weighted) log-likelihood

$$\sum_{y=1}^K \sum_{i=1}^n p(y|i) \log p_y$$

- The resulting estimate (cf. before)

$$\hat{p}_y = \frac{\sum_{i=1}^n p(y|i)}{n}, \quad y = 1, \dots, K$$

M-step: Gaussians

- ▶ We maximize the expected (weighted) log-likelihood (separately for each Gaussian)

$$\sum_{i=1}^n p(y|i) \log N(x^{(i)}; \mu^{(y)}, \sigma_y^2 I)$$

- ▶ The resulting estimates (cf. before)

$$\hat{\mu}^{(y)} = \frac{\sum_{i=1}^n p(y|i) x^{(i)}}{\sum_{i=1}^n p(y|i)}$$

$$\hat{\sigma}_y^2 = \frac{\sum_{i=1}^n p(y|i) \|x^{(i)} - \hat{\mu}^{(y)}\|^2}{d \sum_{i=1}^n p(y|i)}$$



What did we learn?

► **Modeling**

- generative models are specified by variables and how the variables are related
- “latent variables” allow us to specify different underlying structures that we can (try to) uncover

► **Estimation (general)**

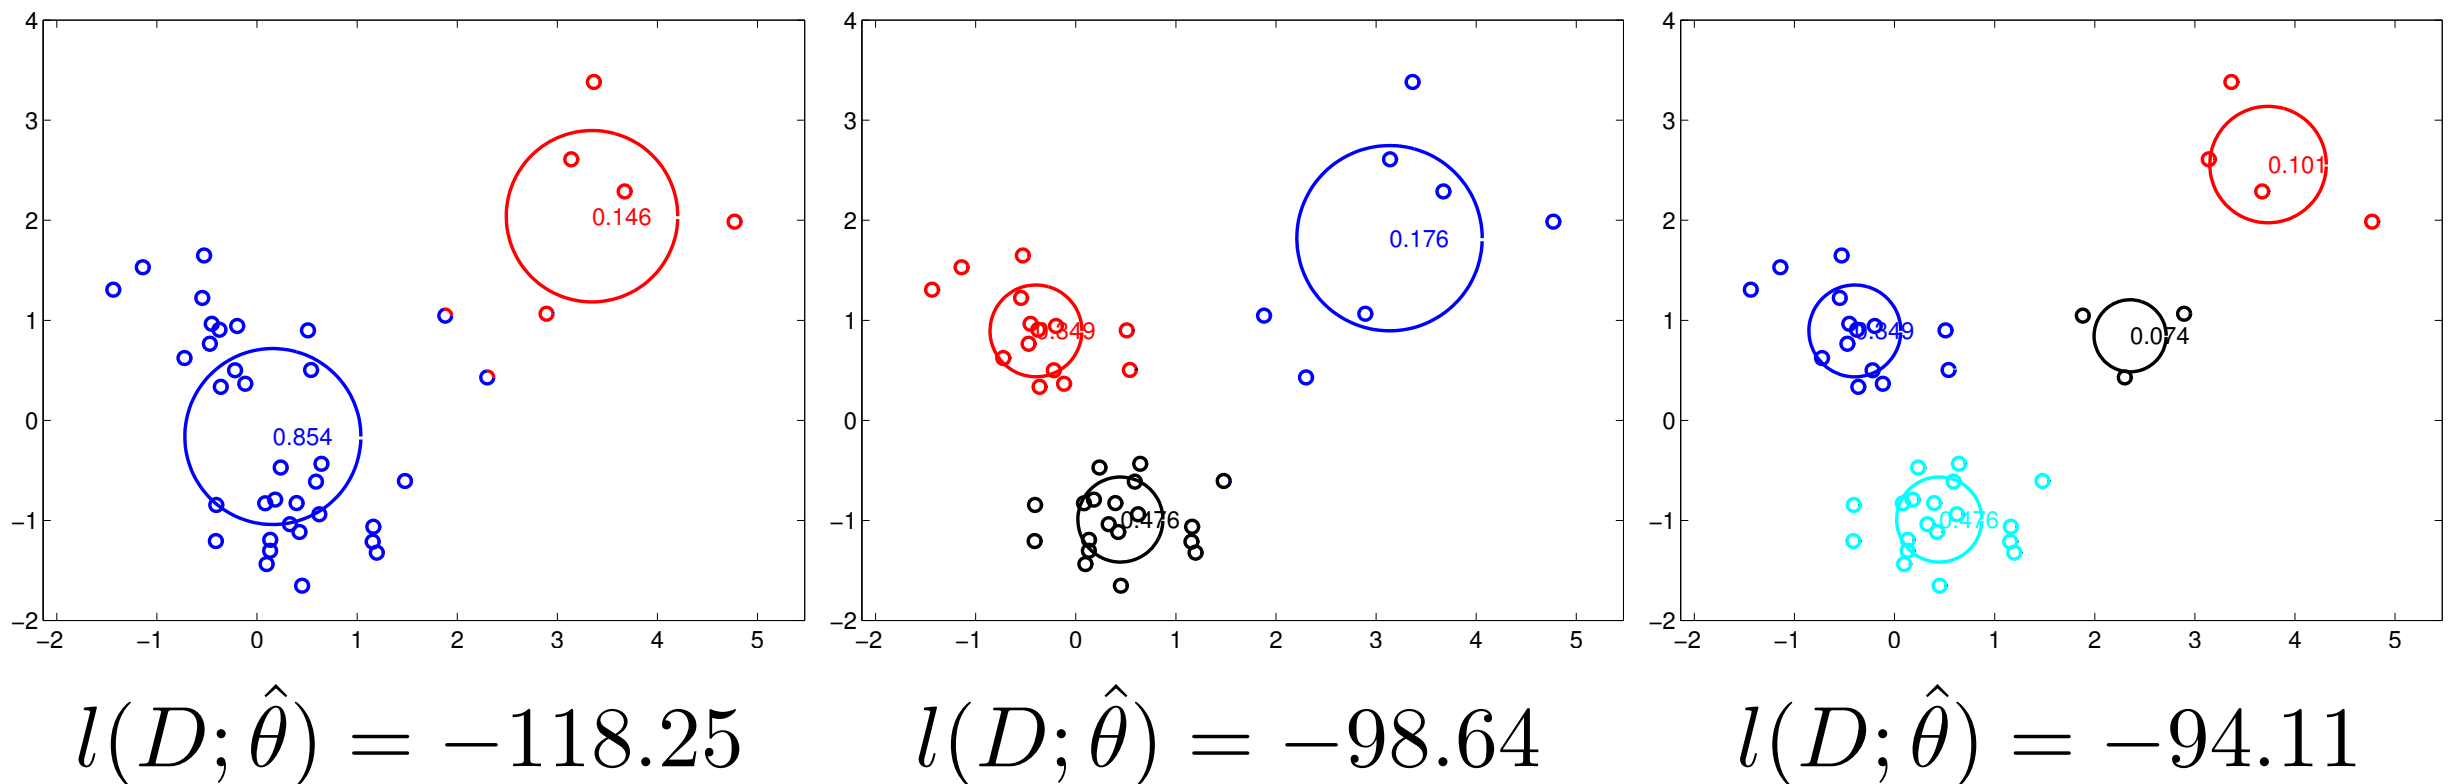
- models can be estimated “in pieces” if we have complete data, i.e., observations with a value assignment for each variable
- easy estimation extends to weighted complete data

► **Estimation (EM algorithm)**

- EM algorithm iteratively creates weighted training sets (based on posterior assignments) and updates the model parameters, separately for each piece, based on the weighted data
- the EM algorithm applies the same to other generative models where y specifies different latent choices (more later)

Model selection

- We can run the EM-algorithm with different numbers of components. Need to specify a criterion for selecting among the different models.



- Basing the selecting on the value of log-likelihood would invariably lead to the largest number of components

Model selection: BIC

- Bayesian Information Criterion (BIC)

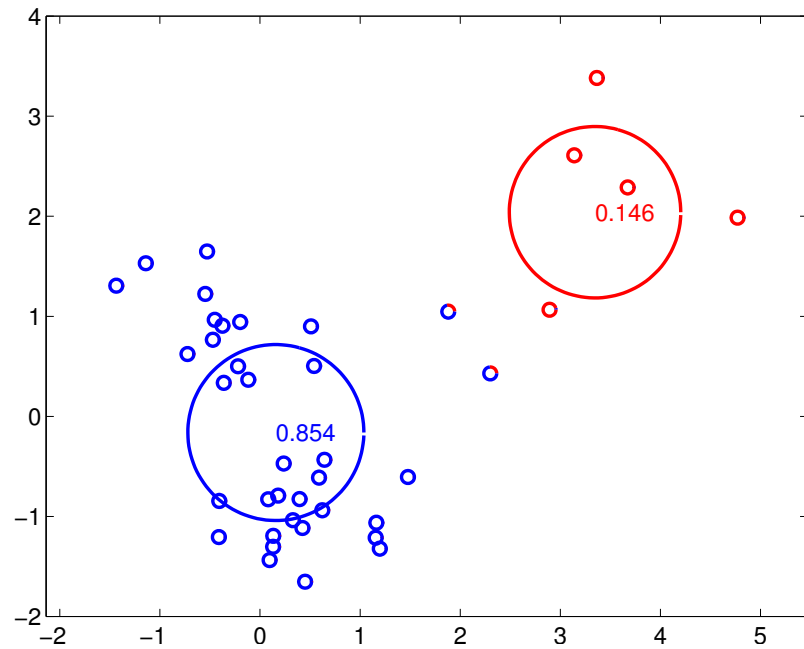
$$BIC(D; \hat{\theta}) = l(D; \hat{\theta}) - \frac{\dim(\hat{\theta})}{2} \log(n)$$

Diagram illustrating the components of the BIC formula:

- $l(D; \hat{\theta})$: ML parameter estimates
- $\frac{\dim(\hat{\theta})}{2}$: # of parameters in the model (complexity penalty)
- $\log(n)$: # of training examples

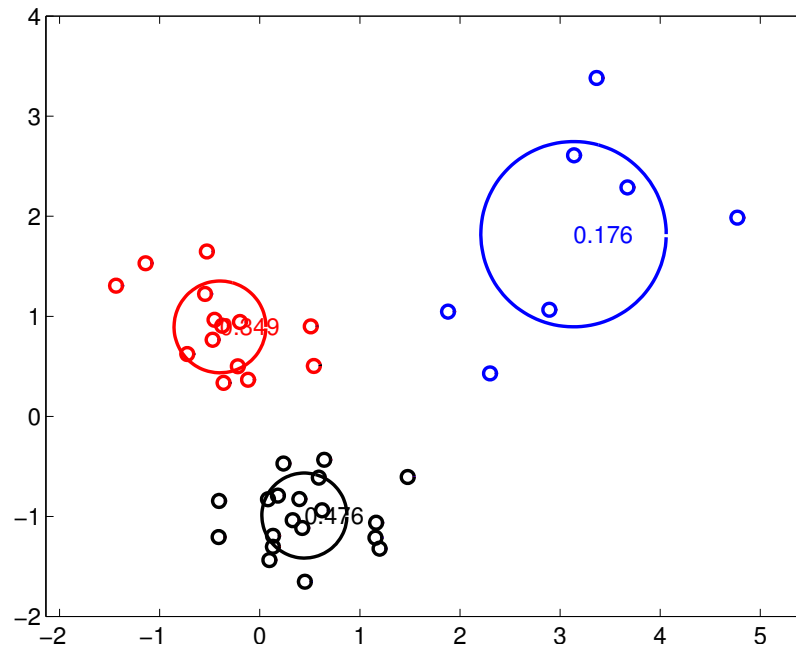
- Many closely related ways to think about “complexity” penalties in model selection criteria
 - how excess parameters help fit the data even in the absence of any signal
 - the cost of communicating the model (likelihood tells us the cost of transmitting the data given the model)
 - fraction of the parameter space that explains the data

Mixture models and BIC



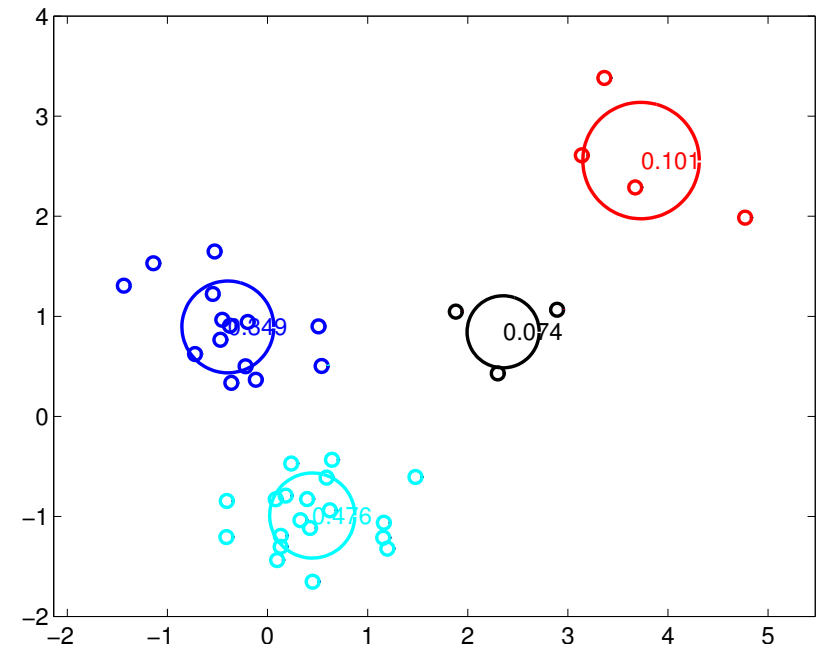
$$l(D; \hat{\theta}) = -118.25$$

$$BIC(D; \hat{\theta}) = -131.16$$



$$l(D; \hat{\theta}) = -98.64$$

$$BIC(D; \hat{\theta}) = -118.93$$



$$l(D; \hat{\theta}) = -94.11$$

$$BIC(D; \hat{\theta}) = -121.78$$



Modeling sequences

- ▶ Lots of interesting data come in the form of sequences
 - temporal data (financial, monitoring, speech)
 - languages
 - user behavior
 - bio-sequences
 - etc.
- ▶ Our goal is to model such sequences, i.e., specify and learn probability distributions over sequences
 - specification (how to define, parameterize)
 - sampling (understand as a generative model)
 - estimation (learn from data)



Modeling sequences

- ▶ Consider a sequence of (e.g., binary) variables

$$y_1 \quad y_2 \quad y_3 \quad y_4 \quad \dots \quad y_i \in V$$

- ▶ We wish to specify a probability distribution over their values. By the chain rule, we can always write

$$P(y_1, \dots, y_n) = P(y_1)P(y_2|y_1)P(y_3|y_2, y_1) \cdots P(y_n|y_{n-1}, \dots, y_1)$$

but, without any assumptions, we would need very large probability tables $|V|^n$

- ▶ The 1st order Markov assumption (bigram model):

$$P(y_1, \dots, y_n) \stackrel{def}{=} P(y_1)P(y_2|y_1)P(y_3|y_2) \cdots P(y_n|y_{n-1})$$

Markov Model

- A 1st order (homogenous) Markov Model requires us to specify two sets of distributions

$$y_1 \rightarrow y_2 \rightarrow \cdots \rightarrow y_{n-1} \rightarrow y_n$$

- Initial state distribution:

$$P(y_1 = y) = \pi_y, \quad \sum_{y=1}^K \pi_y = 1$$

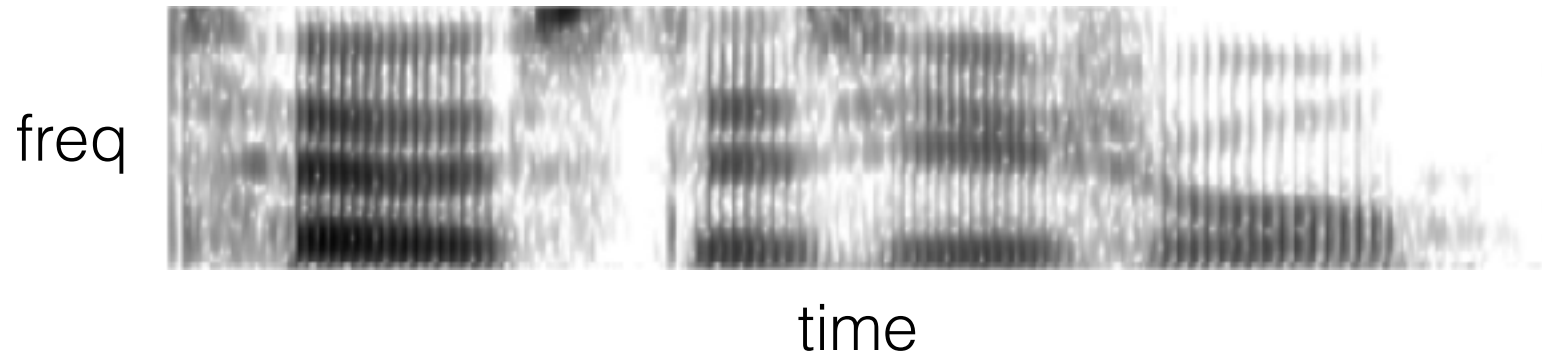
- State transition probabilities:

$$P(y_i = y' | y_j = y) = a_{y,y'}, \quad \sum_{y'=1}^K a_{y,y'} = 1 \quad \forall y$$



Modeling hidden sequences

- Speech recognition



- Handwriting recognition

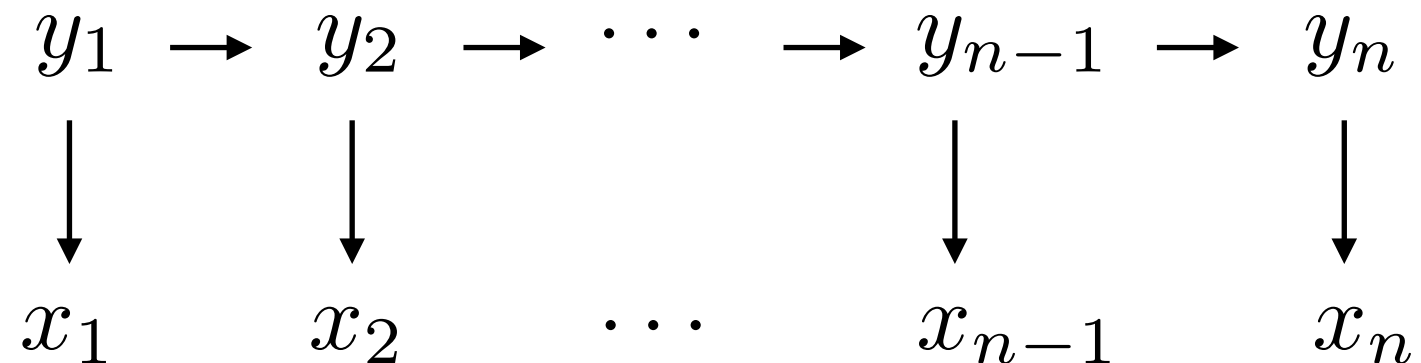
Machine learning is fun

- Information extraction, part of speech tagging, etc.
- Bio-sequence annotation
- Etc.



Modeling hidden sequences

- Similarly to mixture models, we must connect the latent selections to actual observations

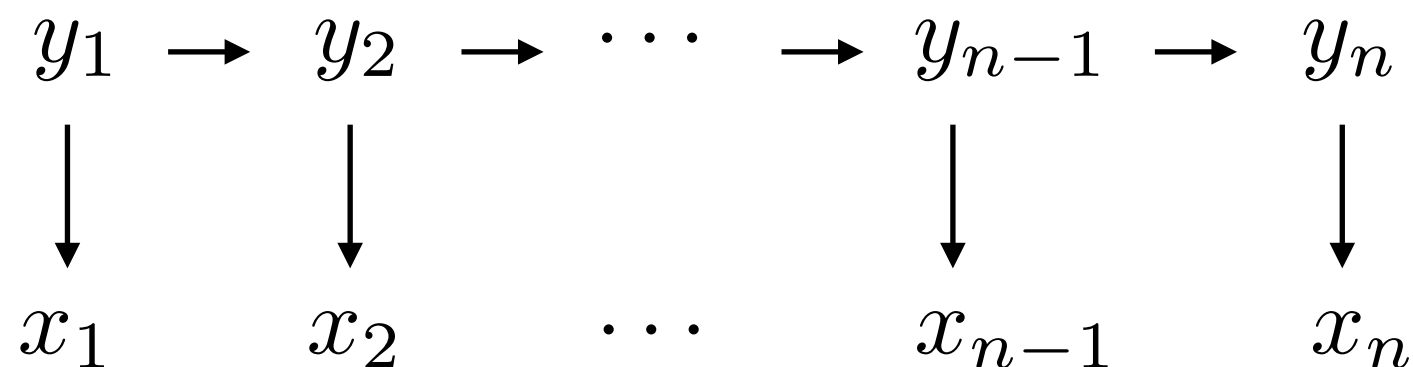


- If the latent sequence is Markov, and the observations are only tied to the state at the corresponding time, we get a Hidden Markov Model (HMM)

$$P(y_1, \dots, y_n, x_1, \dots, x_n) = P(y_1)P(x_1|y_1) \prod_{i=2}^n [P(y_i|y_{i-1})P(x_i|y_i)]$$

HMM

- ▶ We need three sets of probabilities to specify a homogeneous HMM



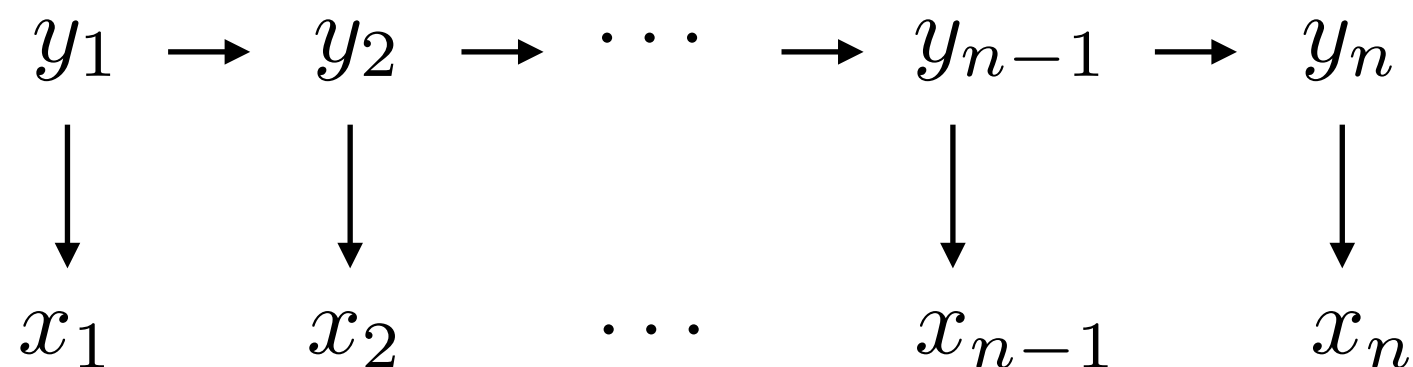
$$P(y_1 = y) = \pi_y, \quad \sum_{y=1}^K \pi_y = 1 \quad \text{initial state distribution}$$

$$P(y_i = y' | y_{i-1} = y) = a_{y,y'}, \quad \sum_{y'=1}^K a_{y,y'} = 1 \quad \forall y \quad \text{state transition}$$

$$P(x_i = x | y_i = y) = \begin{cases} N(x; \mu^{(y)}, \sigma_y^2 I) \\ \text{or discrete} \end{cases} = b_y(x) \quad \text{observation model}$$

HMM

- ▶ We need three sets of probabilities to specify a homogeneous HMM



$$P(y_1 = y) = \pi_y, \quad \sum_{y=1}^K \pi_y = 1 \quad \text{initial state distribution}$$

$$P(y_i = y' | y_{i-1} = y) = a_{y,y'}, \quad \sum_{y'=1}^K a_{y,y'} = 1 \quad \forall y \quad \text{state transition}$$

$$P(x_i = x | y_i = y) = \begin{cases} N(x; \mu^{(y)}, \sigma_y^2 I) \\ \text{or discrete} \end{cases} = b_y(x) \quad \text{observation model}$$

$$P(y_1, \dots, y_n, x_1, \dots, x_n; \theta) = \pi_{y_1} b_{y_1}(x_1) \prod_{i=2}^n [a_{y_{i-1}, y_i} b_{y_i}(x_i)]$$



HMM problems to solve

- ▶ **Evaluation:** calculate the probability of observed values, i.e., sum over all the underlying state sequences

$$P(x_1, \dots, x_n) = \sum_{y_1, \dots, y_n} P(y_1) P(x_1 | y_1) \prod_{i=2}^n [P(y_i | y_{i-1}) P(x_i | y_i)]$$

- ▶ **Estimation:** EM algorithm for HMMs
 - E-step now needs a little effort...
 - M-step is easy (separate estimation of pieces)
- ✓ ▶ **Prediction:** find the most likely hidden state sequence responsible for the observations (Viterbi)

$$\hat{y}_1, \dots, \hat{y}_n = \arg \max_{y_1, \dots, y_n} P(y_1, \dots, y_n, x_1, \dots, x_n)$$

Viterbi algorithm

- ▶ We can find the most likely hidden sequence using dynamic programming (focusing first on the max value)

$$d_i(y_i) = \max_{y_1, \dots, y_{i-1}} P(y_1, \dots, y_i, x_1, \dots, x_i)$$

- ▶ Recursion (should be implemented on a log-scale)

$$d_1(y_1) = P(y_1)P(x_1|y_1)$$

$$d_i(y_i) = \max_{y_{i-1}} \left\{ d_{i-1}(y_{i-1})P(y_i|y_{i-1})P(x_i|y_i) \right\}$$

- ▶ The end result of the recursion is that

$$\max_{y_n} d_n(y_n) = \max_{y_1, \dots, y_n} P(y_1, \dots, y_n, x_1, \dots, x_n)$$



Viterbi, backtracking

- ▶ We can reconstitute the maximizing sequence by working backwards, instantiating the argmax's

$$\hat{y}_n = \arg \max_{y_n} d_n(y_n)$$

$$\hat{y}_{i-1} = \arg \max_{y_{i-1}} \left\{ d_{i-1}(y_{i-1}) P(\hat{y}_i | y_{i-1}) P(x_i | y_i) \right\}$$



Key things to know

- Mixture models
 - specification, sampling, estimation
 - the EM algorithm (for mixtures)
- Markov Models
 - specification, sampling, ML estimation
- Hidden Markov Models
 - specification, sampling, (estimation qualitatively)
 - Viterbi algorithm for prediction